



# Amazon DynamoDB: What It Is and 10 Things You Should Know



## What is Amazon DynamoDB?

Amazon DynamoDB is a managed NoSQL service with strong consistency and predictable performance that shields users from the complexities of manual setup.

### NoSQL vs SQL Databases

Whether or not you've actually used a NoSQL data store yourself, it's probably a good idea to make sure you fully understand the key design differences between NoSQL (including Amazon DynamoDB) and the more traditional relational database (or "SQL") systems like MySQL.

First of all, **NoSQL** does not stand for "**Not SQL**", but "**Not Only SQL**". The two are not opposites, but complementary. NoSQL designs deliver faster data operations and can seem more intuitive, while not necessarily adhering to the ACID (atomicity, consistency, isolation, and durability) properties of a relational database.

There are many well-known NoSQL databases available, including MongoDB, Cassandra, HBase, Redis, [Amazon DynamoDB](#), and Riak. Each of these was built for a specific range of uses and will therefore offer different features. We could group these databases into the following categories: *columnar* (Cassandra, HBase), *key-value store* (DynamoDB, Riak), *document-store* (MongoDB, CouchDB), and *graph* (Neo4j, OrientDB).

In this post, I'm going to focus on AWS DynamoDB — the giant of the NoSQL world. I believe it's become a giant because AWS built it for their own operations. Considering how much was at stake financially, anything less than complete reliability would simply not be tolerated. Software created in such a demanding environment and with the use of AWS-scale resources is bound to be epic. The result? Fantastic reliability and durability, with blazing fast service.

Like any other [AWS](#) product, Amazon DynamoDB was designed for failure (i.e., it has self-recovery and resilience built in). That makes DynamoDB a highly available, scalable, and distributed data store. Here are ten DynamoDB features that helped make this database service into a giant.

## 1. Amazon DynamoDB is a managed, NoSQL database service

With a managed service, users only interact with the running application itself. You don't need to worry about things like server health, storage, and network connectivity. With Amazon DynamoDB, AWS provisions and runs the infrastructure for you. Some of DynamoDB's critical managed infrastructure features include:

- Automatic data replication over three availability-zones in a single region.
- Infinitely scalable read-write I/O running on IOPS-optimized solid state drives.
- A provisioned-throughput model where read and write units can be adjusted at any time based on actual application usage.
- Data backed up to S3.
- Integrated with other AWS services like Elastic MapReduce ([Amazon EMR](#)), Data Pipeline, and [Amazon Kinesis](#).
- Amazon DynamoDB pricing: Pay-per-use model – you never pay for hardware or services you're not actually using.
- Security and access control can be applied using [AWS IAM](#) service.
- Great enterprise-ready features such as a robust SLA, monitoring tools, and private VPN functionality.

## 2. Amazon DynamoDB has Predictable Performance

AWS claims that DynamoDB will deliver highly predictable performance. Considering Amazon's reputation for service delivery, we tend to take them at their word on this one. You can actually control the quality of the service you'll get by choosing between Strong Consistency (Read-after-Write) or Eventual Consistency. Similarly, if a user wants to increase or decrease the Read/Write throughput they'll experience, they can do it through simple API calls. Amazon DynamoDB also offers what they call **Provisioned Capacity**: briefly, DynamoDB allows you to "bank" up to five minutes of unused capacity, which, like the funds in an emergency bank account, you can use during short bursts of activity.

## 3. Amazon DynamoDB is designed for massive scalability

Being an AWS product, you can assume that Amazon DynamoDB will be extremely scalable. With their automatic partitioning model, as data volumes grow, DynamoDB spreads the data across partitions and raises throughput. This requires no intervention from the user.

## 4. Amazon DynamoDB data types

DynamoDB supports the following data types:

- Scalar** – Number, String, Binary, Boolean, and Null.
- Multi-valued** – String Set, Number Set, and Binary Set.
- Document** – List and Map.

Scalar types are generally well understood. We'll focus instead on multi-valued and document types. Multi-valued types are sets, which means that the values in this data type are unique. For a *months* attribute you can choose a String Set with the names of all twelve months – each of which is, of course, unique.

Similarly, document types are meant for representing complex data structures in the form of Lists and Maps. See this example:

```
{
  Id = 100
  ProductName = "K3 Note"
```

```
color = { "white", "black" }
ProductCategory = "Mobile"
}
```

## 5. Amazon DynamoDB’s Data Model

DynamoDB uses three basic data model units, Tables, Items, and Attributes. Tables are collections of Items, and Items are collections of Attributes.

Amazon DynamoDB attributes are basic units of information, like key-value pairs. Tables are like tables in relational databases, except that in DynamoDB, tables do not have fixed schemas associated with them. Items are like rows in an RDBMS table, except that DynamoDB requires a Primary Key. The Primary Key in DynamoDB must be unique so that it can find the exact item in the table. DynamoDB keys supports two kinds of Primary Keys:

- Hash Type Primary Key:** If an attribute uniquely identifies an item, it can be considered as Primary. DynamoDB builds a hash index on the attribute to facilitate the uniqueness. A Hash Key is mandatory in a DynamoDB table.
- Hash and Range Type Primary Key:** This type of Primary Key is built upon the hashed key and the range key in the table: a hashed index on the hash primary key attribute, and a range sort index on the range primary key attribute. This type of primary key allows for AWS’s rich query capabilities.

## 6. Amazon DynamoDB indexes

There are two types of indexes in DynamoDB, a Local Secondary Index (LSI) and a Global Secondary Index (GSI). In an LSI, a range key is mandatory, while for a GSI you can have either a hash key or a hash+range key. GSIs span multiple partitions and are placed in separate tables. DynamoDB supports up to five GSIs. While creating a GSI, you need to carefully choose your hash key because that key will be used for partitioning.

Which is the right index type to use? Here are two considerations: LSIs limit item size to 10 GB, and GSIs offer only eventual consistency.

## 7. Amazon DynamoDB partitions

In DynamoDB, data is partitioned automatically by its hash key. That’s why you will need to choose a hash key if you’re implementing a GSI. The partitioning logic depends upon two things: table size and throughput.



The partition for a table is calculated by DynamoDB. Although it is transparent to users, you should understand the logic behind this.

(**Note:** Read Capacity Units – RCU – are measured in 4KB/sec. Write Capacity Units – WCU – are measured in 1KB/sec.)  
According to this formula, if we have a table size of 16 GB and we have 6000 RCUs and 1000 WCUs, then:

# of partitions by throughput:  $6000/3000+1000/1000 = 3$

# of partitions by size:  $16/10 = 1.6$

So, the # of partitions in total:  $\max(1.6, 3) = 3$

Therefore, we will require three partitions. The RCUs and WCUs will be uniformly distributed across the partitions. Here, RCUs per partition will be  $3000/3 = 1000$ . RCUs and the WCUs will be  $1000/3 = 333$  WCUs. The data per partition will be  $16/3 = 5.4$  GB per partitions.

## 8. Amazon DynamoDB streams

DynamoDB streams are like transactional logs for a table. According to the DynamoDB Developer’s Guide:

*A DynamoDB stream is an ordered flow of information about changes to items in an Amazon DynamoDB table. When you enable a stream on a table, DynamoDB captures information about every modification to data items in the table.*

Streams are applied only to tables, and each stream record appears exactly once in a stream. AWS maintains separate endpoints for DynamoDB and DynamoDB streams. There are all kinds of scenarios where streams can be useful, such as in a messaging application where a message or picture that is updated to a group must be reflected in the message boxes of all the group members, Or for sending welcome messages to new customers when they sign up for your service.

## 9. Amazon DynamoDB integration with Amazon EMR and Redshift

NoSQL and Big Data technologies are often discussed together, because they both share the same distributed and horizontally scalable architecture, and both aim to provide high volume, structured, and semi-structured data processing. In a typical scenario, Elastic MapReduce (EMR) performs its complex analysis on datasets stored on DynamoDB. Users will often also use AWS Redshift for data warehousing, where BI tasks are carried out on data loaded from DynamoDB tables to Redshift.

## 10. Amazon DynamoDB JavaScript Web Shell

AWS has introduced a web-based user interface known as the DynamoDB JavaScript Shell for local development. You can download the tool [here](#).

### Steps:

- Download and extract the appropriate file
- Run following command:

```
java -Djava.library.path=./DynamoDBLocal_lib -jar DynamoDBLocal.jar
```

Access the console in a browser with the URL: **http://localhost:8000/shell**

The web page will look like this:

Click on the button to get some sample commands

For example the createTable API will run:

After running this, listTables will show you:

This is a great tool to perform syntax checking before actually going to production.

With DynamoDB, Amazon has done a great job providing a NoSQL service with strong consistency and predictable performance, while saving users from the complexities of a distributed system. One proof of their success is the many systems (like Riak) that chose to build on the DynamoDB design. With a strong ecosystem, Amazon DynamoDB is something to consider when you are building your next Internet-based scale application.

Ready to try it for yourself? Why not use Cloud Academy's [AWS DynamoDB hands-on lab](#)?

AWS DynamoDB



WRITTEN BY CHANDAN PATRA

Cloud Computing and Big Data professional with 10 years of experience in pre-sales, architecture, design, build and troubleshooting with best engineering practices. Specialities: Cloud Computing - AWS, DevOps(Chef), Hadoop Ecosystem, Storm & Kafka, ELK Stack, NoSQL, Java, Spring, Hibernate, Web Service



### Subscribe to our blog

Get the latest posts in your email

OR [SUBSCRIBE TO THE](#) [RSS FEED](#)

## The cloud skills platform of choice for teams & innovators

Certification Learning Paths.  
Hands-on Labs. Skill Validation.  
Assess your team's tech skills today.

Tour the Product

CARLOS RIVAS

DECEMBER 1, 2023

AWS re:Invent

New AWS re:Invent Announcements: Dr. Werner Vogels Keynote

Read more

ALANA LAYTON

DECEMBER 1, 2023

AWS re:Invent

New AWS re:Invent Announcements: AWS Partner Keynote with Dr Ruba Borno

Read more

DANNY JESSEE

NOVEMBER 30, 2023

AWS re:Invent

New AWS re:Invent Announcements: Swami Sivasubramanian Keynote

Read more

STUART SCOTT

AWS re:Invent

AWS re:Invent announcements Keynote

Read more

### SUGGESTED POSTS

- Cloud Academy's Black Friday Deal Is Here!
- The Black Friday Early-Bird Deal Starts Now!
- Cloud Academy Wins Again: Fall 2023 G2 Awards
- Navigating the Challenges of Modernization with Cloud Academy
- What is Cloud Migration? Strategy, Processes, Benefits & Risks
- Google Cloud Certifications: Which is Right for You and Your Team?
- The 12 AWS Certifications: Which is Right for You and Your Team?
- Level Up Your Cloud Skills: Oracle Now in Cloud Academy's Training Library
- Get 40% Off Individual Plans at Cloud Academy
- Retaining Your Rockstars: How to Avoid Management Disruptions
- 5 Reasons Why Top Consulting Firms Fail at Multi-Cloud – And How to Fix It

Workforce Transformation: Building Tech Talent From Within

Our Black Friday Preview Gets You 20% Off!

Get Back to School with 20% Off Personal Plans

Get 40% off with our Pre-Summer Event!

Skills Intelligence Part 2: The State of Tech Skills Training

Deciding on Your Cloud Certification

Skills Intelligence Part 1: Baseline Your Team's Tech Skills

Cloud Academy's Free Tech Skills Weekend is April 22-25

How to Become a Software Engineer

Real-World Cloud Talk: An Interview with a Microsoft Sr. Cloud Solution Architect

Cloud Academy Remains a Leader in the G2 Spring 2022 Reports

Navigating the Great Resignation

What is Skills Intelligence?

New: AI on Alibaba, Terraform Labs on Google Cloud, plus more

Baseline Skills to Make the Right Hire

5 Tech Skills to Build in 2022

Start Your New Year With Cloud Academy!

New Year's Resolution Sale!

New AWS re:Invent Announcements: Dr. Werner Vogels Keynote

re:Invent 2021: Faster Cloud Migrations with AWS Mainframe Modernization Platform

New AWS re:Invent Announcements: Swami Sivasubramanian Keynote

re:Invent 2021: AWS Announces New Amazon Inspector

New AWS re:Invent Announcements: Adam Selipsky Keynote

The AWS re:Invent Global Partner Summit Keynote: All You Need to Know

AWS Certified Solutions Architect Associate: A Study Guide

AI, GitHub + More – 5 Key Announcements from Microsoft Ignite 2021

An Interview With a Real Cloud Marathoner

The Biggest Challenges for Technology Leaders

Why Skills Development Is Critical for Tech Success

Cloud Migration Series (Step 5 of 5): Manage & Iterate

Join the Cloud Marathon Today

Cloud Migration Series (Step 4 of 5): Adopt a Cloud-First Mindset

It's 10:00 AM: Do You Know Where Your Team's Tech Skills Are?

Cloud Migration Series (Step 3 of 5): Assess Readiness

Cloud Migration Series (Step 2 of 5): Start Planning

Cloud Migration Series (Step 1 of 5): Define Your Strategy

Jump Into Cloud Academy's Tech Skills Assessment

The Positive Side of 2020: People — and Their Tech Skills — Are Everyone's Priority

Give the Gift of Cloud Academy

Cloud Academy Referrals: Get \$20 for Every Friend Who Subscribes!

The Art of the Exam: Get Ready to Pass Any Certification Test

Cloud Skills and Real Guidance for Your Organization: Our Special Campaign Begins!

AWS Certification Practice Exam: What to Expect from Test Questions

Cloud Academy Nominated High Performer in G2 Summer 2020 Reports

AWS Security: Bastion Hosts, NAT instances and VPC Peering

AWS Security Groups: Instance Level Security

What Exactly Is a Cloud Architect and How Do You Become One?

How to Use & Install the AWS CLI

S3 FTP: Build a Reliable and Inexpensive FTP Server Using Amazon's S3

How DNS Works - the Domain Name System (Part One)

What is Ansible?

Google Unveils Gemini AI

New AWS re:Invent Announcements: Dr. Werner Vogels Keynote

New AWS re:Invent Announcements: AWS Partner Keynote with Dr Ruba Borno

New AWS re:Invent Announcements: Swami Sivasubramanian Keynote

AWS re:Invent 2023: New announcements - Adam Selipsky Keynote

EXPLORE LIBRARY →

Pricing About Careers Partners Contact Instructors

SOLUTIONS

Managing the Multi-Cloud

Migrate to the Cloud

Develop Tech Talent

Optimize Cloud Spend

CERTIFICATIONS

Cloud Certifications

AWS Certifications

Azure Certifications

Google Cloud Certifications

RESOURCES

Cloud Webinars

Cloud Conferences

Cloud Case Studies

Whitepapers

Cloud Academy Blog

Press Releases

Courses

Release Notes

TRAINING LIBRARY

AWS Training

Azure Training

Google Cloud Training

Alibaba Cloud Training

Big Data Training

Business Management Training

AWS Certifications

Azure Certifications

GCP Certifications

Alibaba Cloud Certifications

Microsoft 365 Certifications

DevOps Certifications

Security Certifications

Cloud Computing Fundamentals

Cloud Migration Training

Containers Training

DevOps Training

Upcoming Training

PAST EVENTS

COURSE INDEX

Amazon Web Services

Microsoft Azure

Google Cloud Platform

Oracle Cloud Infrastructure

Terms and Conditions Privacy Policy Sitemap System Status

Copyright © 2023 Cloud Academy Inc. All rights reserved.