

# Distributed Systems

## CS 15-440

Consistency and Replication – Part II

Lecture 11, Oct 10, 2011

Majd F. Sakr, **Vinay Kolar**, Mohammad Hammoud

# Today...

- Last Session
  - Consistency and Replication
    - Introduction and Data-Centric Consistency Models
- Today's session
  - Consistency and Replication – Part II
    - Finish Data-centric Consistency Models
    - Client-Centric Consistency Models
    - Replica Management
- Announcement:
  - Interim design report for Project 2 due today

# Recap: Trade-offs in Maintaining Consistency

- + Maintaining consistency should balance between the strictness of consistency versus efficiency
  - + How much consistency is “good-enough” depends on the application

**Loose Consistency**



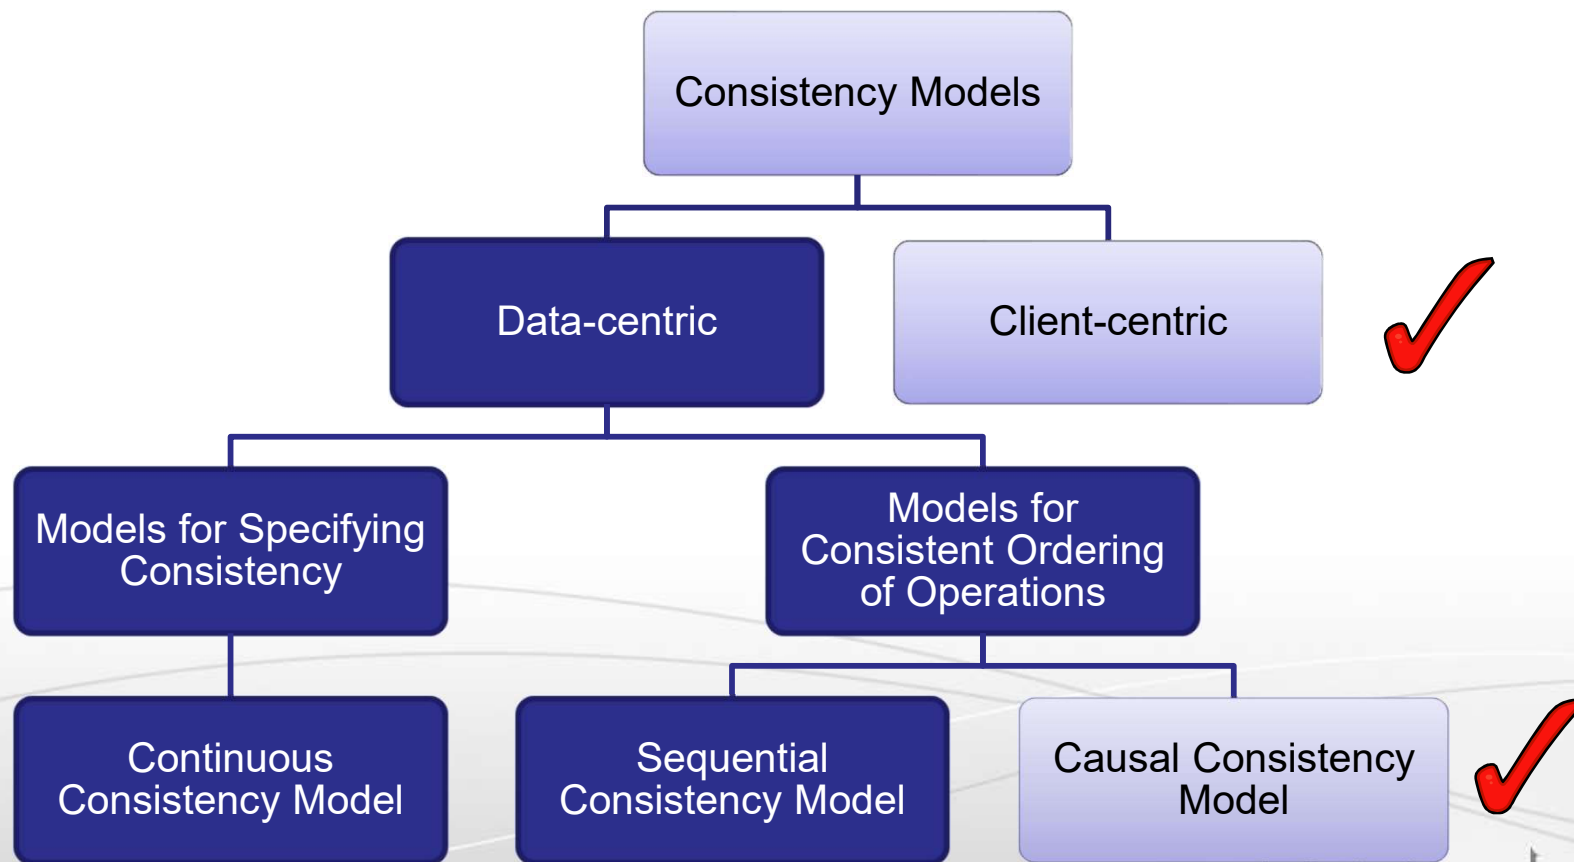
**Strict Consistency**

Easier to implement,  
and is efficient

Generally hard to implement,  
and is inefficient

# Recap: Consistency Models

- + A consistency model states the level of consistency provided by the *data-store* to the processes while reading and writing the data



# Types of Ordering


1. Total Ordering
2. Sequential Ordering
3. Causal Ordering

# Causality (Recap)

- + Causal relation between two events
  - + If **a** and **b** are two events **a** and **b** such that **a** happened-before **b** or  $a \rightarrow b$ , and
  - + If the (logical) time when event **a** and **b** is received at a process  $P_i$  is denoted by  $C_i(a)$  and  $C_i(b)$
  - + Then, if we can infer that  $a \rightarrow b$  by observing that  $C_i(a) < C_i(b)$ , then **a** and **b** are causally related
- + Causality can be implemented using Vector Clocks

# Causal vs. Concurrent events


- + Consider an interaction between processes  $P_1$  and  $P_2$  operating on replicated data  $x$  and  $y$



Events are causally related

Events are not concurrent

- Computation of  $y$  at  $P_2$  may have depended on value of  $x$  written by  $P_1$



Events are not causally related

Events are concurrent

- Computation of  $y$  at  $P_2$  does not depend on value of  $x$  written by  $P_1$

$P_1$  = Process  $P_1$      $\rightarrow$  = Timeline at  $P_1$

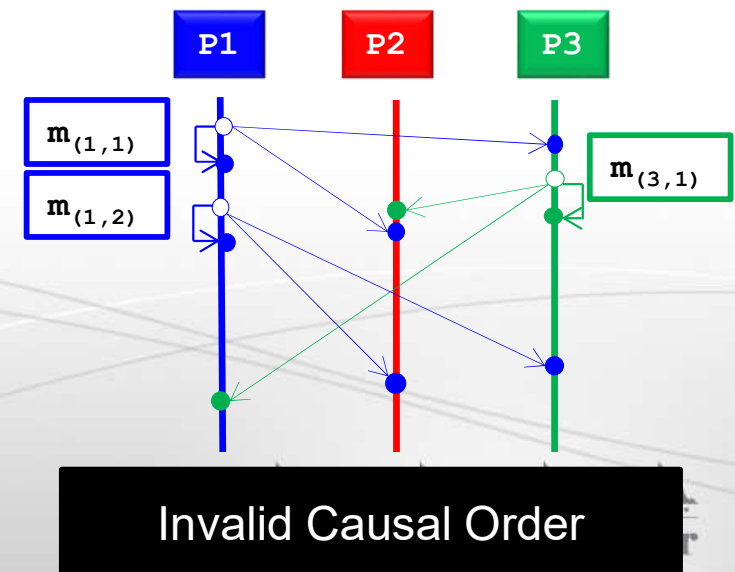

$R(x) b$  = Read variable  $x$ ;  
Result is  $b$

$W(x) b$  = Write variable  $x$ ;  
Result is  $b$

# Causal Ordering

## + Causal Order

- + If process  $P_i$  sends a message  $m_i$  and  $P_j$  sends  $m_j$ , and if  $m_i \rightarrow m_j$  (operator ' $\rightarrow$ ' is Lamport's **happened-before** relation) then any correct process that delivers  $m_j$  will deliver  $m_i$  before  $m_j$
- + In the example,  $m_{(1,1)}$  and  $m_{(3,1)}$  are in Causal Order
- + Drawback:
  - + The **happened-before** relation between  $m_i$  and  $m_j$  should be induced before communication

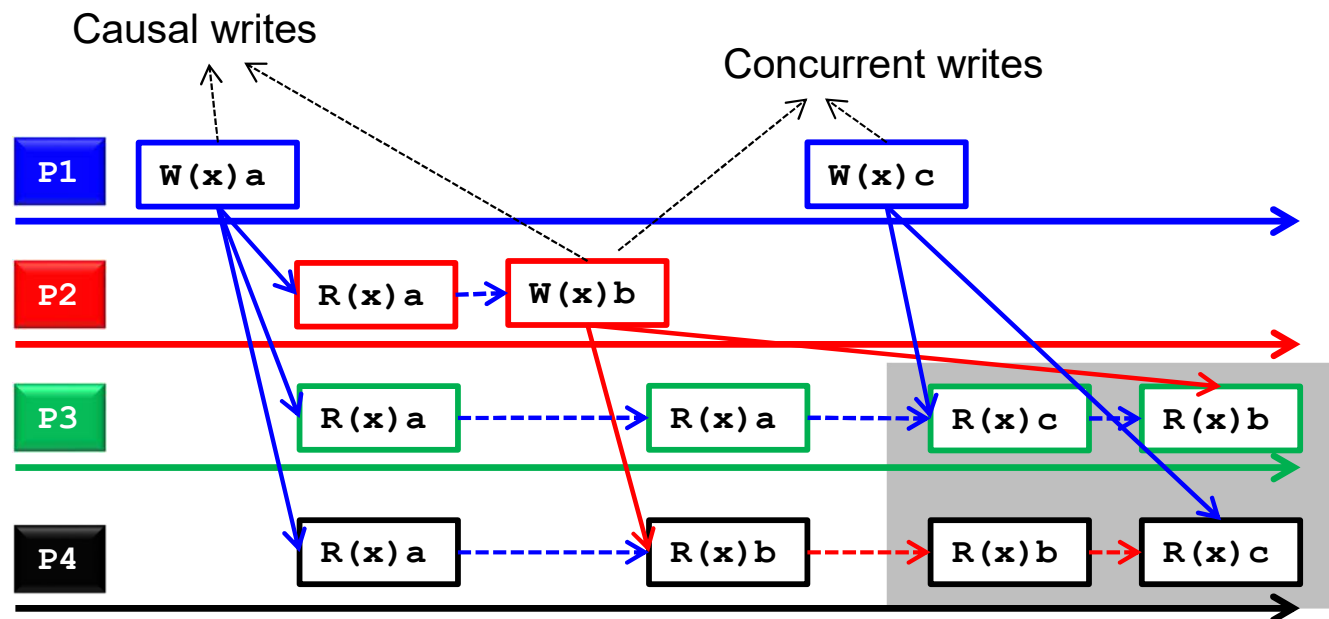




# Causal Consistency Model

- + A data-store is causally consistent if:
  - + Writes that are potentially causally related must be seen by all the processes in the same order
  - + Concurrent writes may be seen in a different order on different machines

# Example of a Causally Consistent Data-store



A Causally Consistent  
Data-Store

But not a Sequentially  
Consistent Data-Store

P1

=Process P1



=Timeline at P1

$R(x) b$

=Read variable  $x$ ;  
Result is  $b$

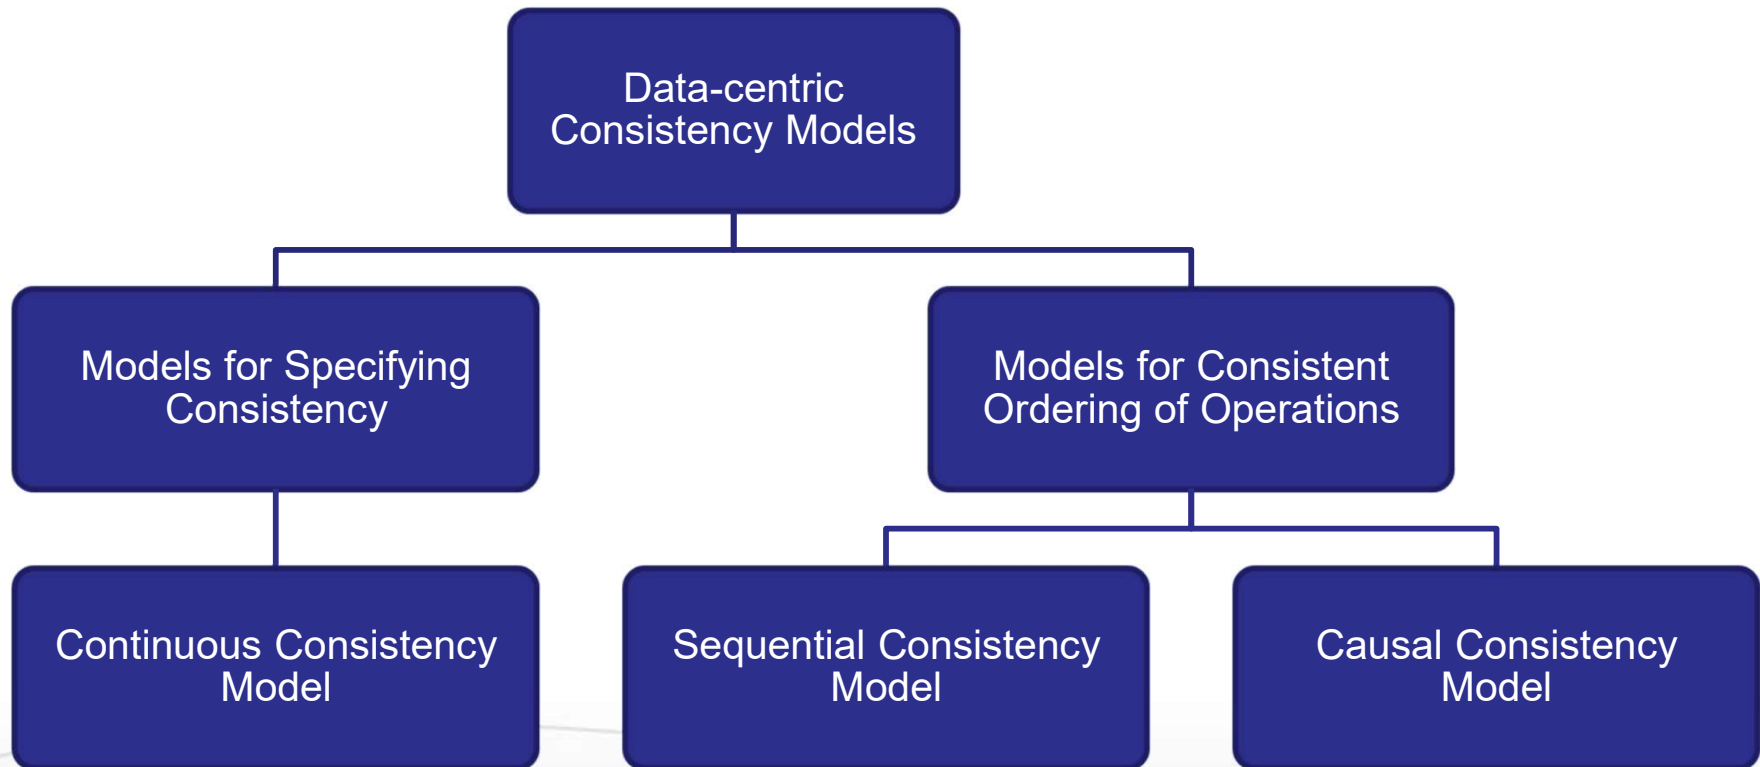
$W(x) b$

= Write variable  $x$ ;  
Result is  $b$

# Implications of adopting a Causally Consistent Data-store for Applications

- + Processes have to keep track of which processes have seen which writes
- + This requires maintaining a dependency graph between write and read operations
  - + Vector clocks provides a way to maintain causally consistent data-base

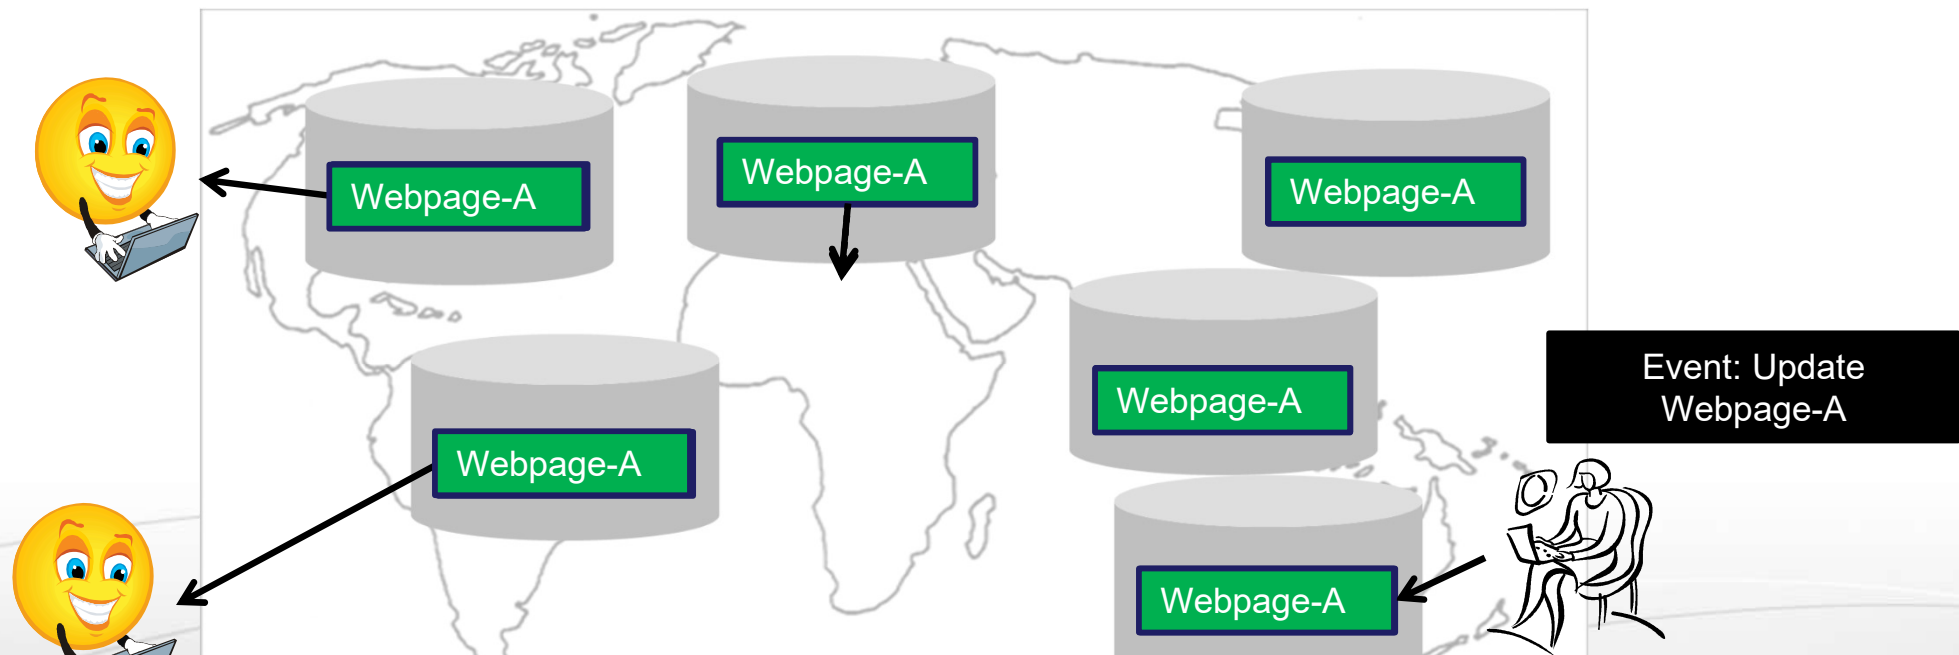
# Topics Covered in Data-centric Consistency Models



But, is Data-centric Consistency Model good for all applications?

# Applications that can use Data-centric Models

- + Data-centric models are applicable when many processes are concurrently updating the data-store
- + But, do all applications need all replicas to be consistent?



Data-Centric Consistency Model is too strict when

- One client process updates the data
- Other processes read the data, and are OK with reasonably stale data

# Summary of Data-Centric Consistency Models

Data-centric consistency models describe how the replicated data is kept consistent across different data-stores, and what the process can expect from the data-store

These models allow measuring and specifying the consistency levels that are tolerable to the application

## Data-centric Consistency Models

These models specify what ordering of operations are ensured at the replicas

### Models for Specifying Consistency

#### Continuous Consistency Model

### Models for Consistent Ordering of Operations

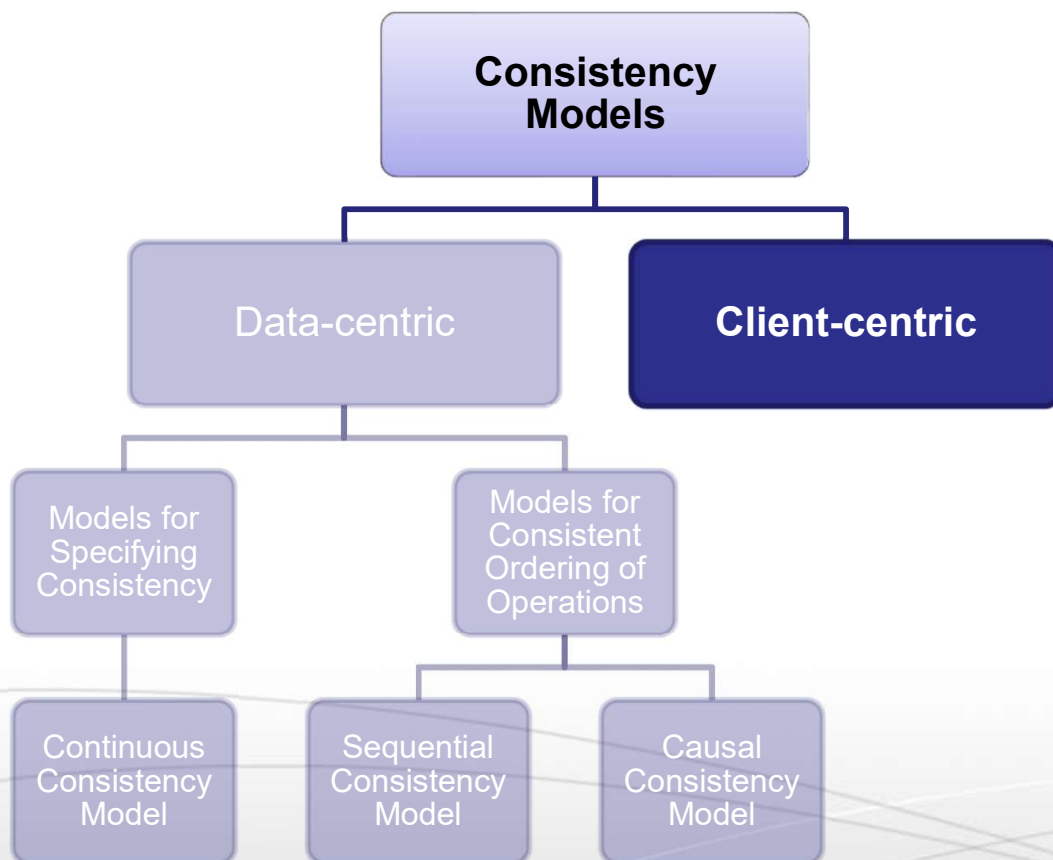
#### Sequential Consistency Model

#### Causal Consistency Model

Data-centric models are too strict when:

- most operations are read operations
- updates are generally triggered from one client process

# Overview

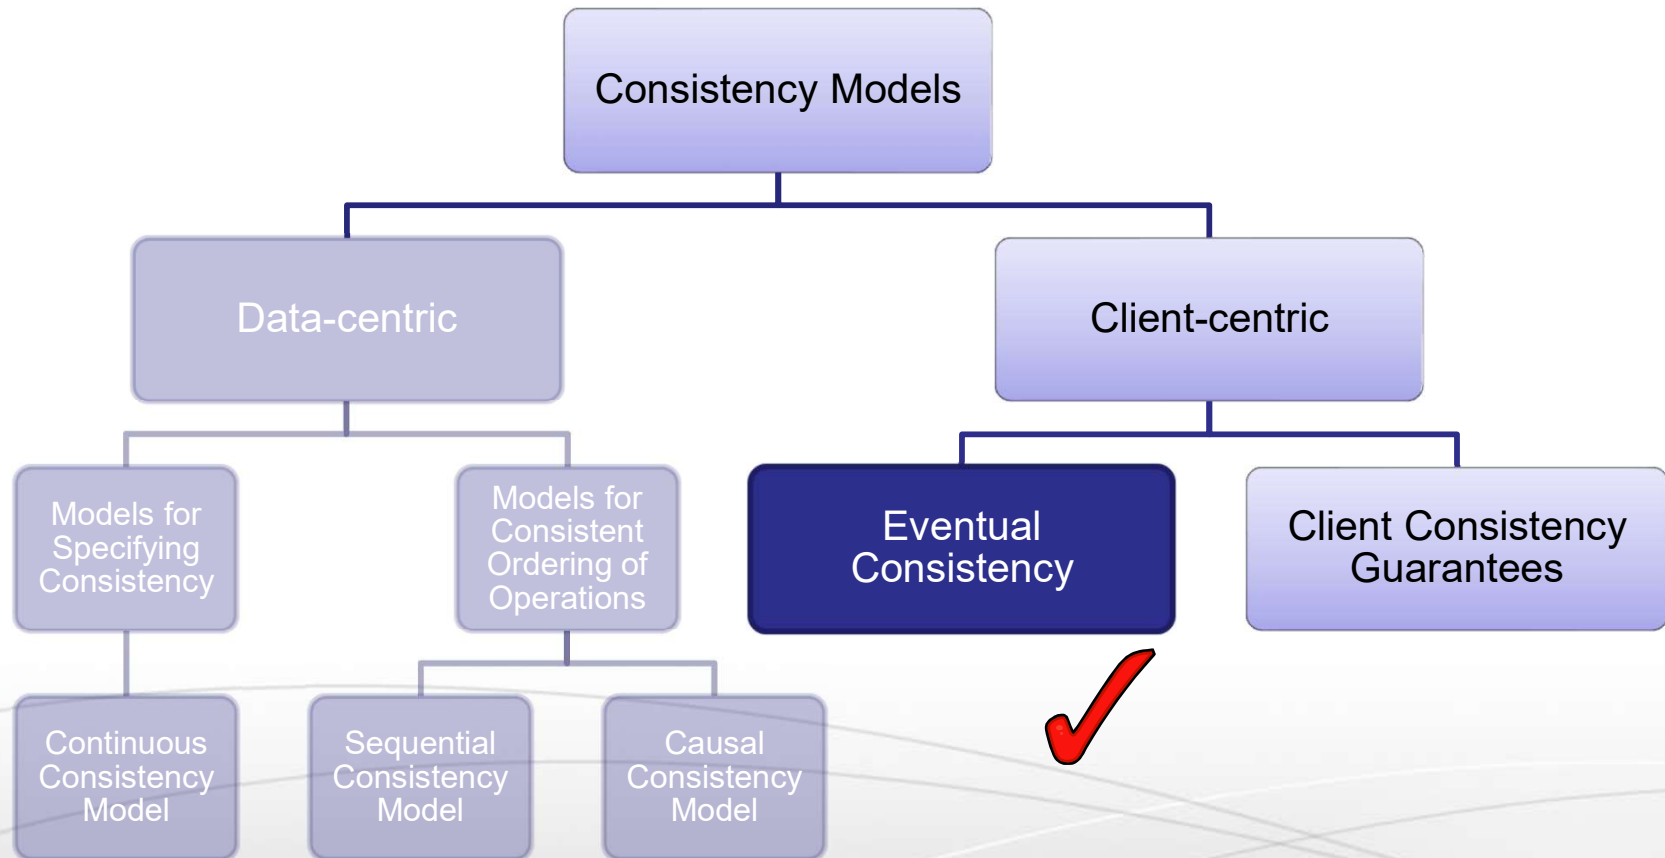


# Client-Centric Consistency Models

- + Data-centric models lead to excessive overheads in applications where:
  - + a majority operations are reads, and
  - + updates occur frequently, and are often from one client process
- + For such applications, a weaker form of consistency called *Client-centric Consistency* is employed for improving efficiency
- + Client-centric consistency models specify two requirements:
  1. **Eventual Consistency**
    - + All the replicas should *eventually* converge on a final value
  2. **Client Consistency Guarantees**
    - + Each client processes should be guaranteed some level of consistency while accessing the data value from different replicas



# Overview



# Eventual Consistency

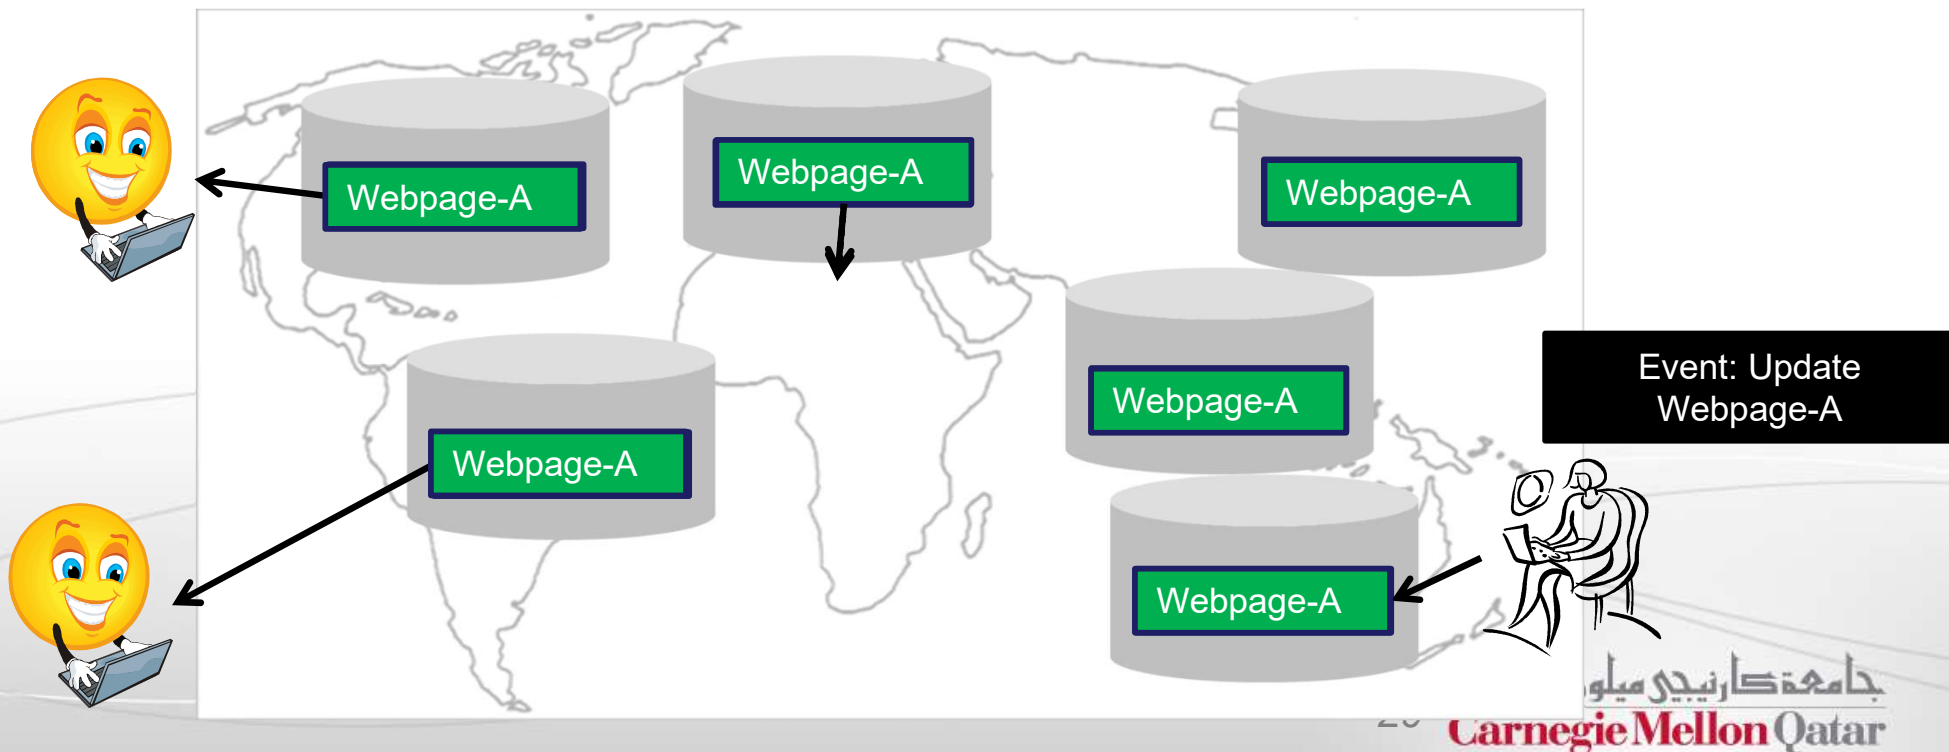
- + Many applications can tolerate a inconsistency for a long time
  - + Webpage updates, Web Search – Crawling, indexing and ranking, Updates to DNS Server
- + In such applications, it is acceptable and efficient if replicas in the data-store rarely exchange updates
- + A data-store is termed as *Eventually Consistent* if:
  - + All replicas will gradually become consistent in the absence of updates
- + Typically, updates are propagated infrequently in eventually consistent data-stores

# Designing Eventual Consistency

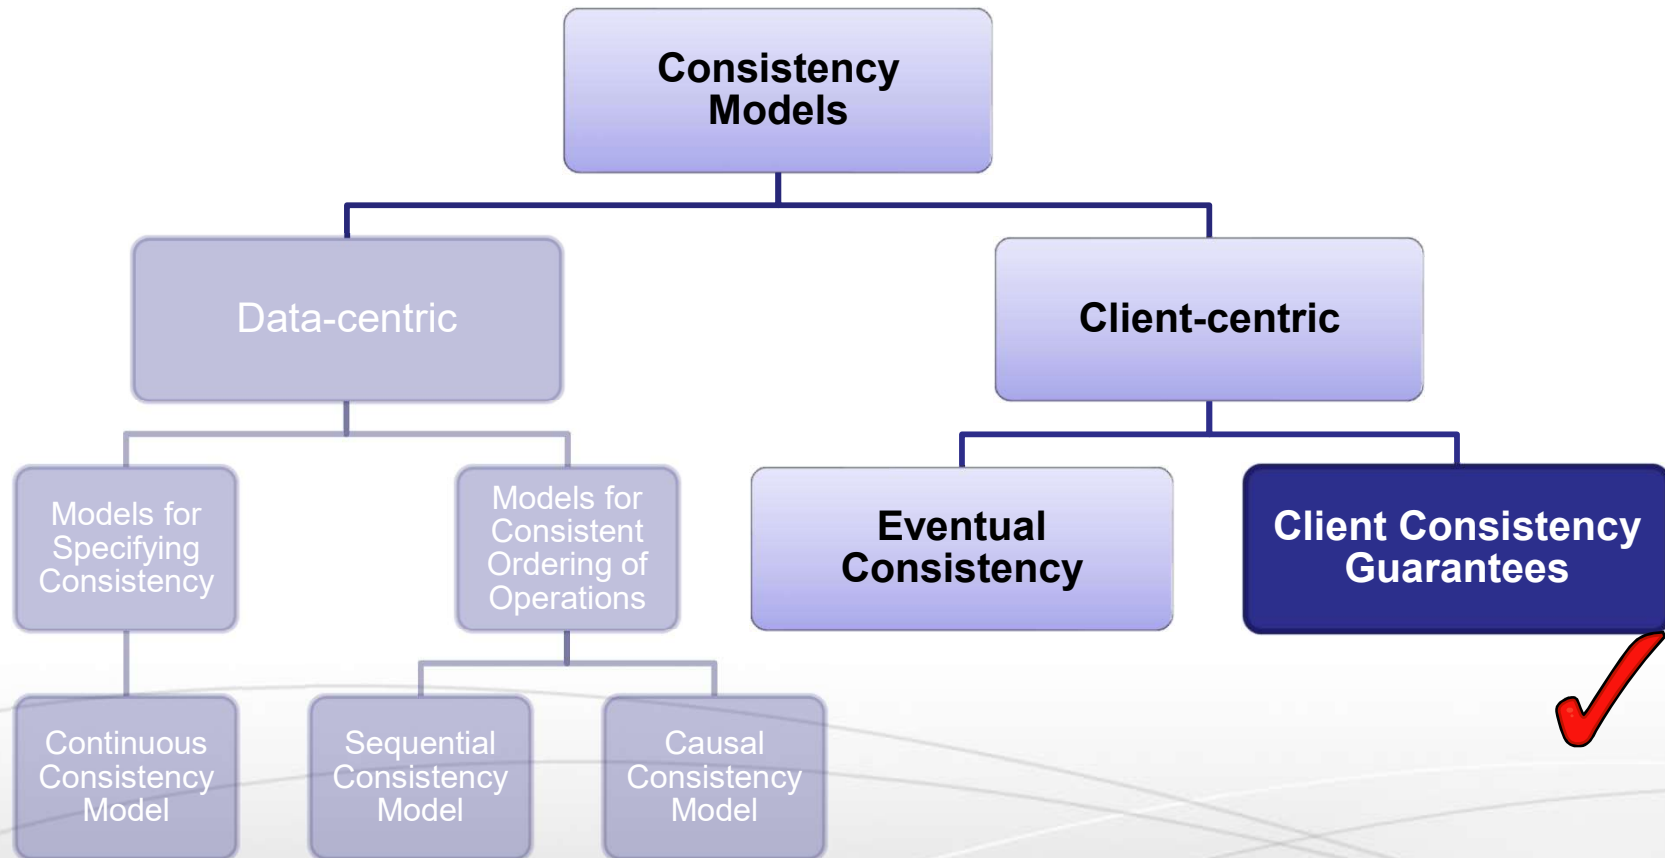
- + In eventually consistent data-stores,
  - + *Write-write conflicts* are rare
    - + Two processes that write the same value are rare
    - + Generally, one client updates the data value
      - + e.g., One DNS server updates the name to IP mapping
    - + Such rare conflicts can be handled through simple mechanisms, such as mutual exclusion
  - + *Read-write conflict* are more frequent
    - + Conflicts where one process is reading a value, while another process is writing a value to the same variable
    - + Eventual Consistency Design has to focus on efficiently resolving such conflicts

# Challenges in Eventual Consistency

- + Eventual Consistency is not good-enough when the client process accesses data from different replicas
  - + We need consistency guarantees for a single client while accessing the data-store




# Overview



# Client Consistency Guarantees

- + Client-centric consistency provides guarantees for a single client for its accesses to a data-store
- + Example: Providing consistency guarantee to a client process for data  $x$  replicated on two replicas. Let  $x_i$  be the local copy of a data  $x$  at replica  $L_i$ .



**$WS(x_1)$**  = Write Set for  $x_1$  = Series of ops being done at some replica that reflects how  $L_1$  updated  $x_1$  till this time

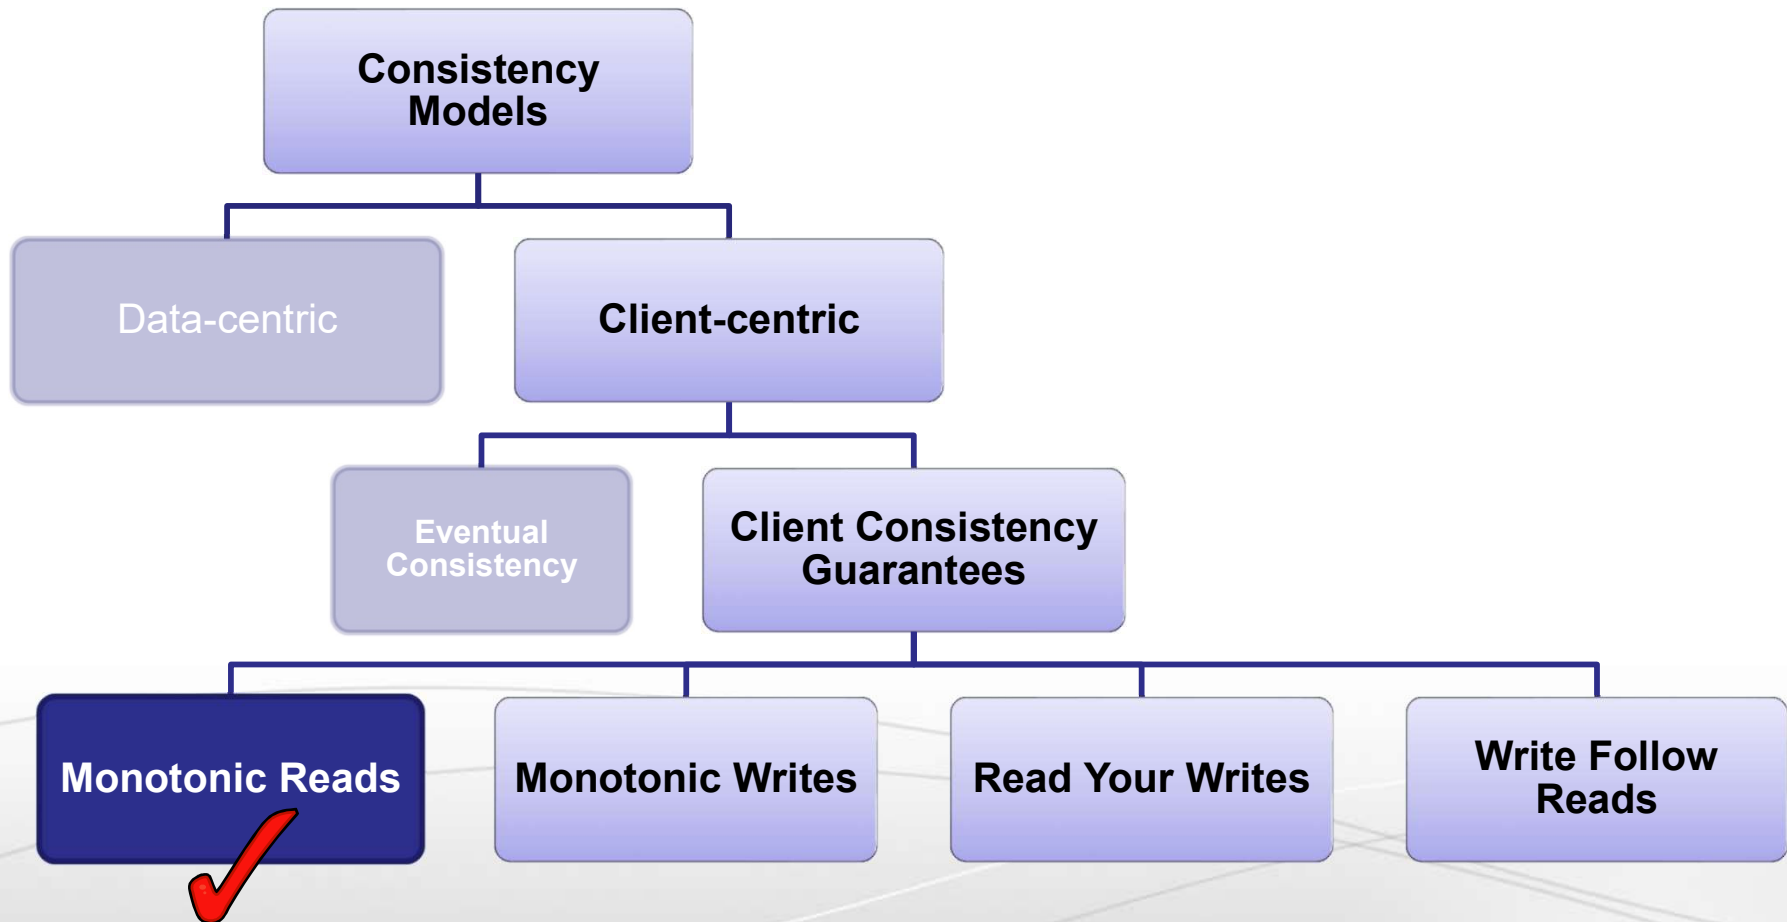
**$WS(x_1; x_2)$**  = Write Set for  $x_1$  and  $x_2$  = Series of ops being done at some replica that reflects how  $L_1$  updated  $x_1$  and, later on, how  $x_2$  is updated on  $L_2$

**$L_i$**  = Replica  $i$        **$R(x_i) b$**  = Read variable  $x$  at replica  $i$ ; Result is  $b$        **$W(x) b$**  = Write variable  $x$  at replica  $i$ ; Result is  $b$        **$WS(x_i)$**  = Write Set

# Client Consistency Guarantees

- + We will study four types of client-centric consistency models<sup>1</sup>
  1. Monotonic Reads
  2. Monotonic Writes
  3. Read Your Writes
  4. Write Follow Reads

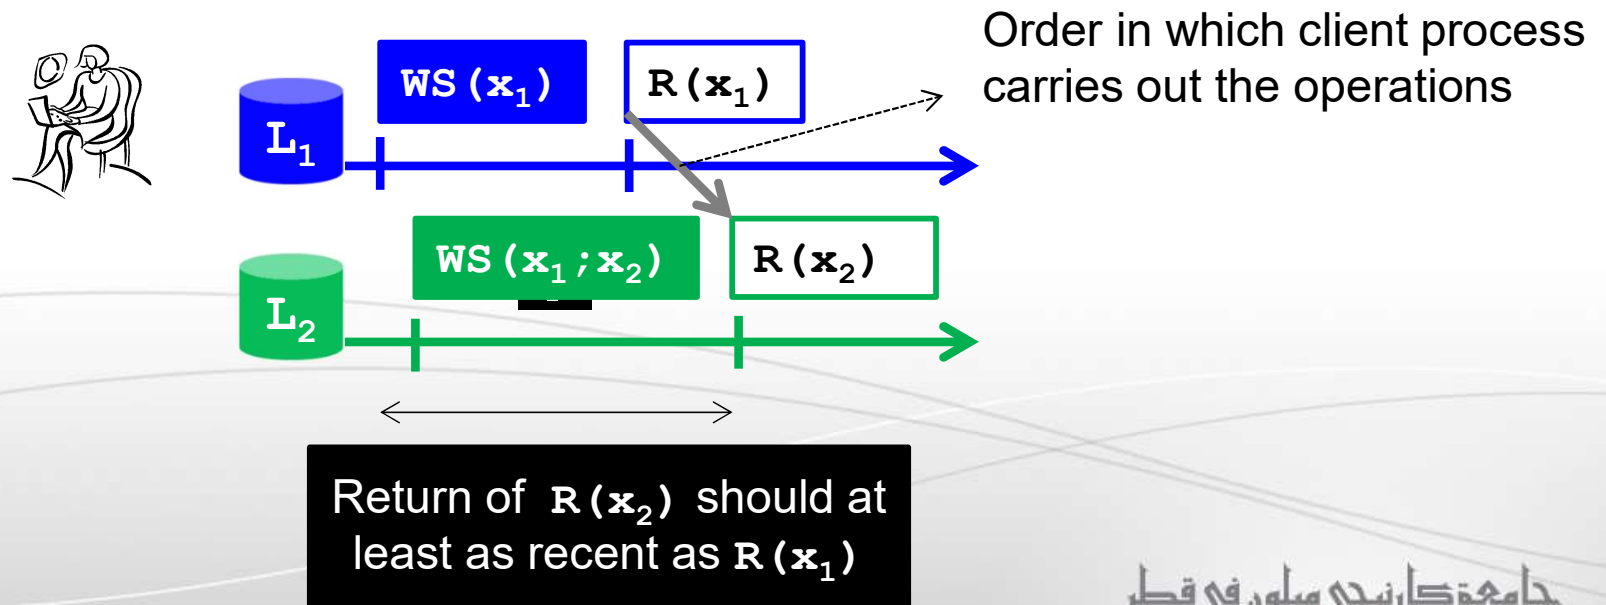
# Overview





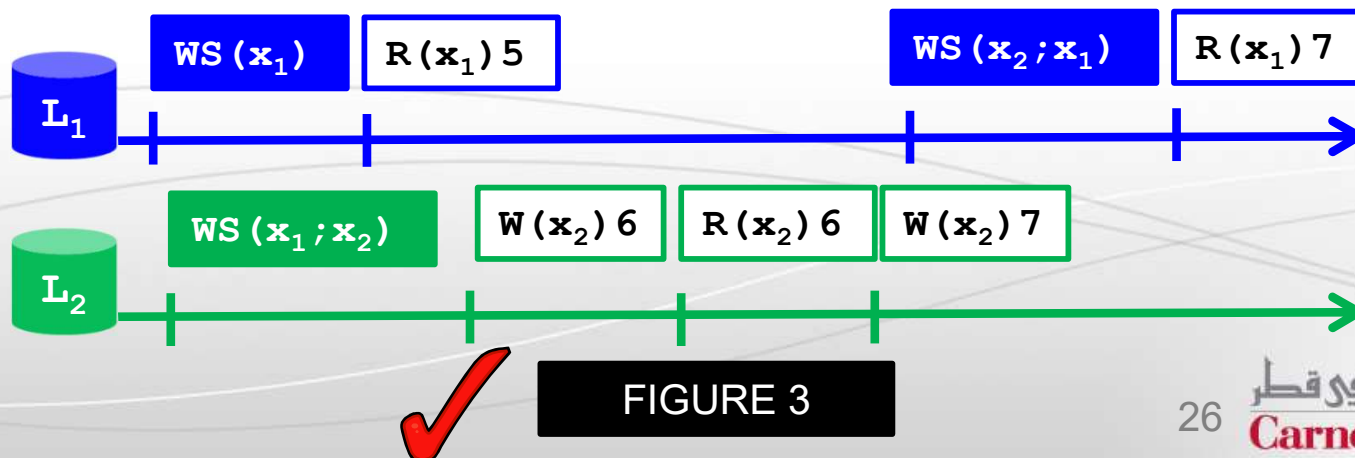
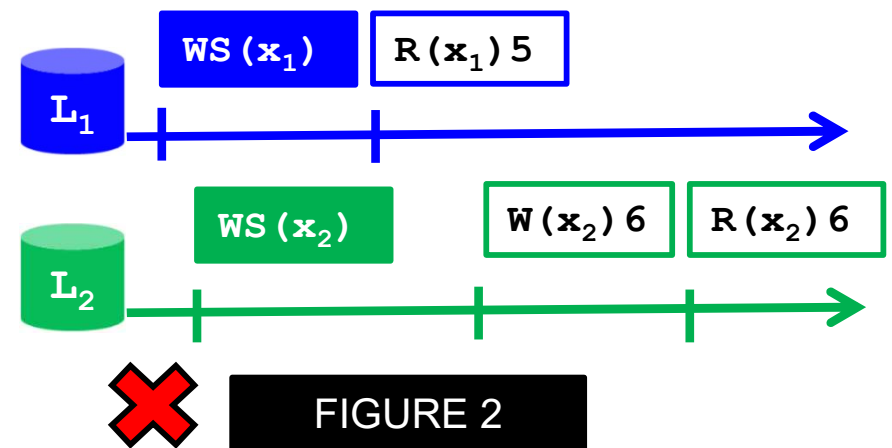
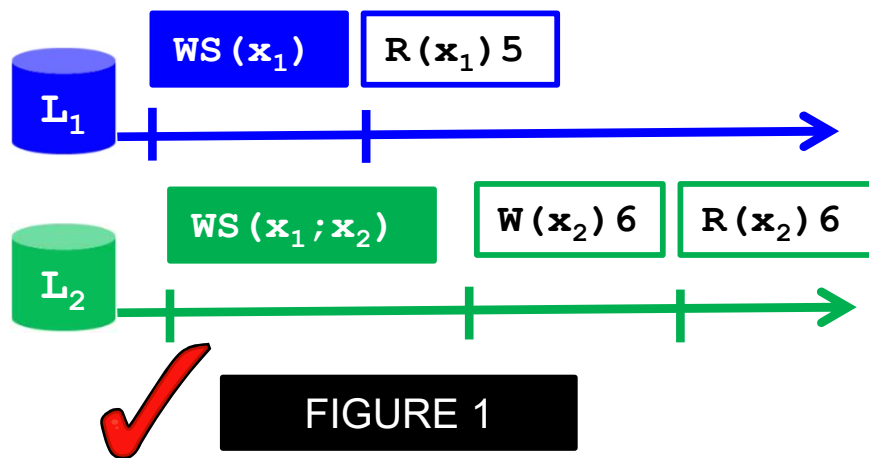
# Monotonic Reads

- + The model provides guarantees on successive reads
- + If a client process reads the value of data item  $x$ , then any successive read operation by that process should return the same or a more recent value for  $x$

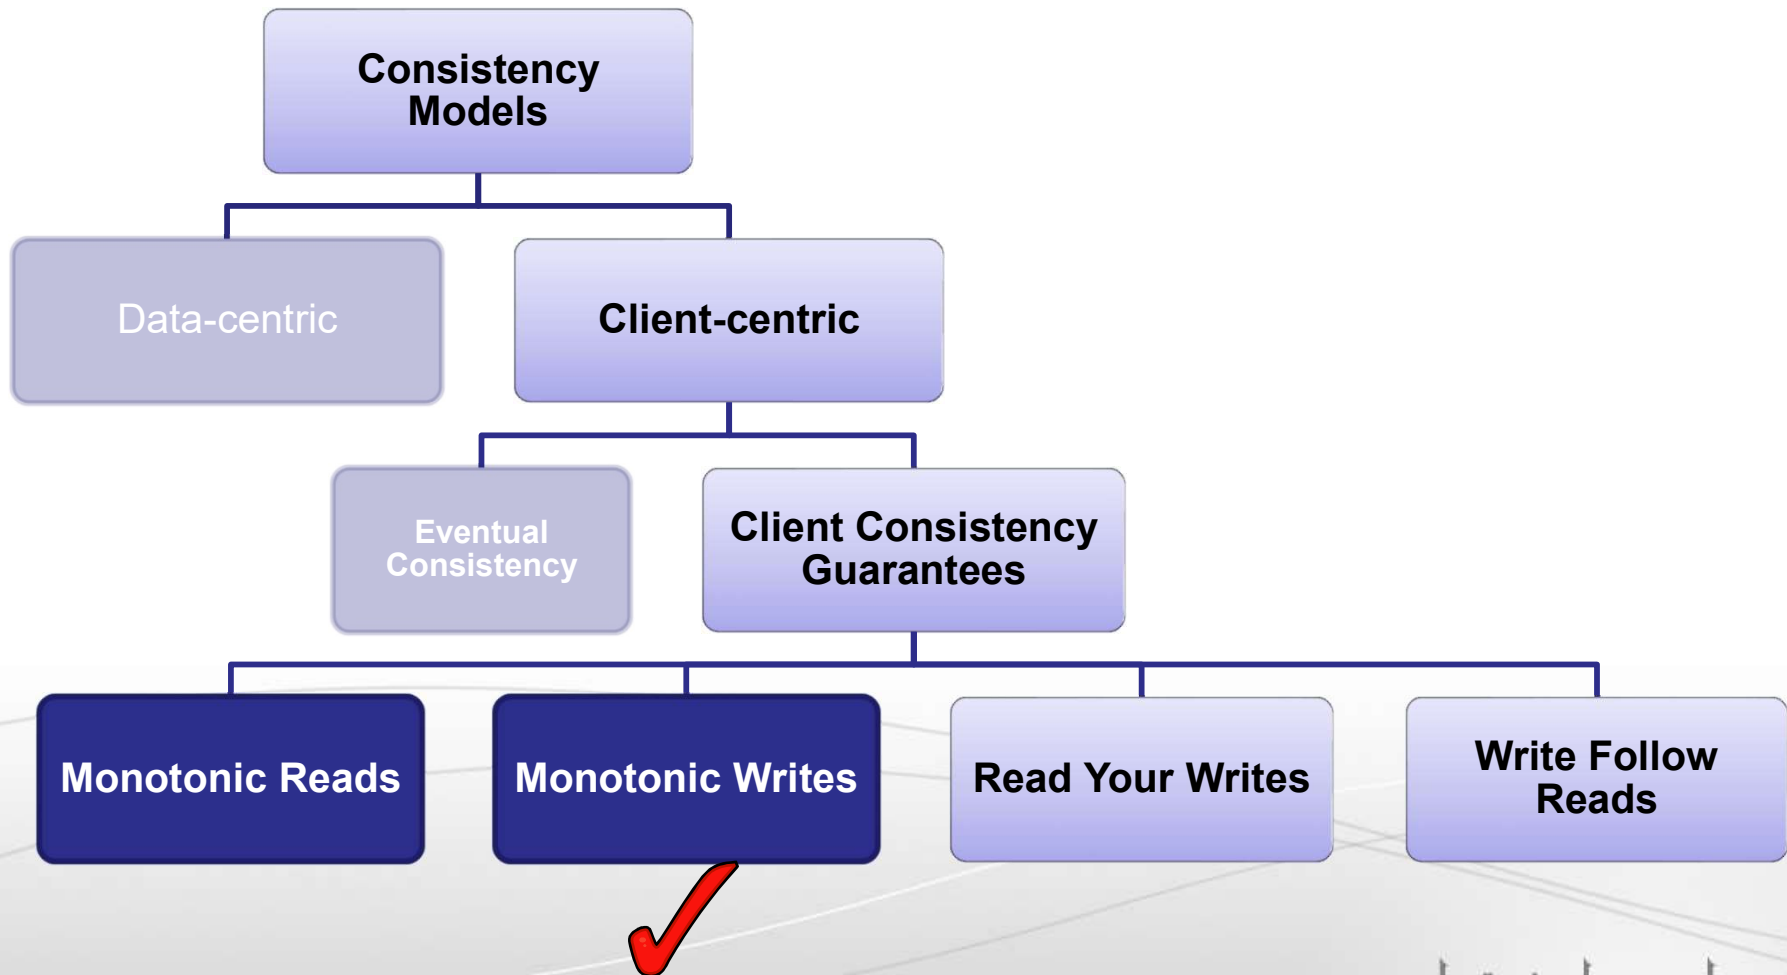


# Monotonic Reads – Puzzle

Recognize data-stores that provide monotonic read guarantees

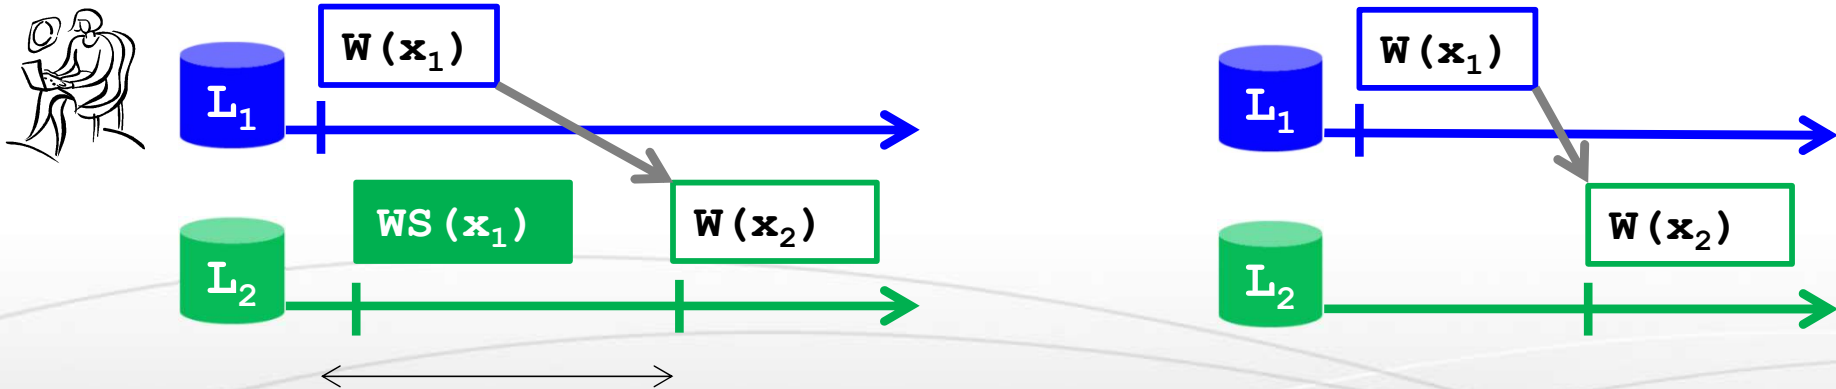


# Overview



# Monotonic Writes

- + This consistency model assures that writes are monotonic
- + A write operation by a client process on a data item  $x$  is completed before any successive write operation on  $x$  by the same process
  - + A new write on a replica should wait for all old writes on any replica



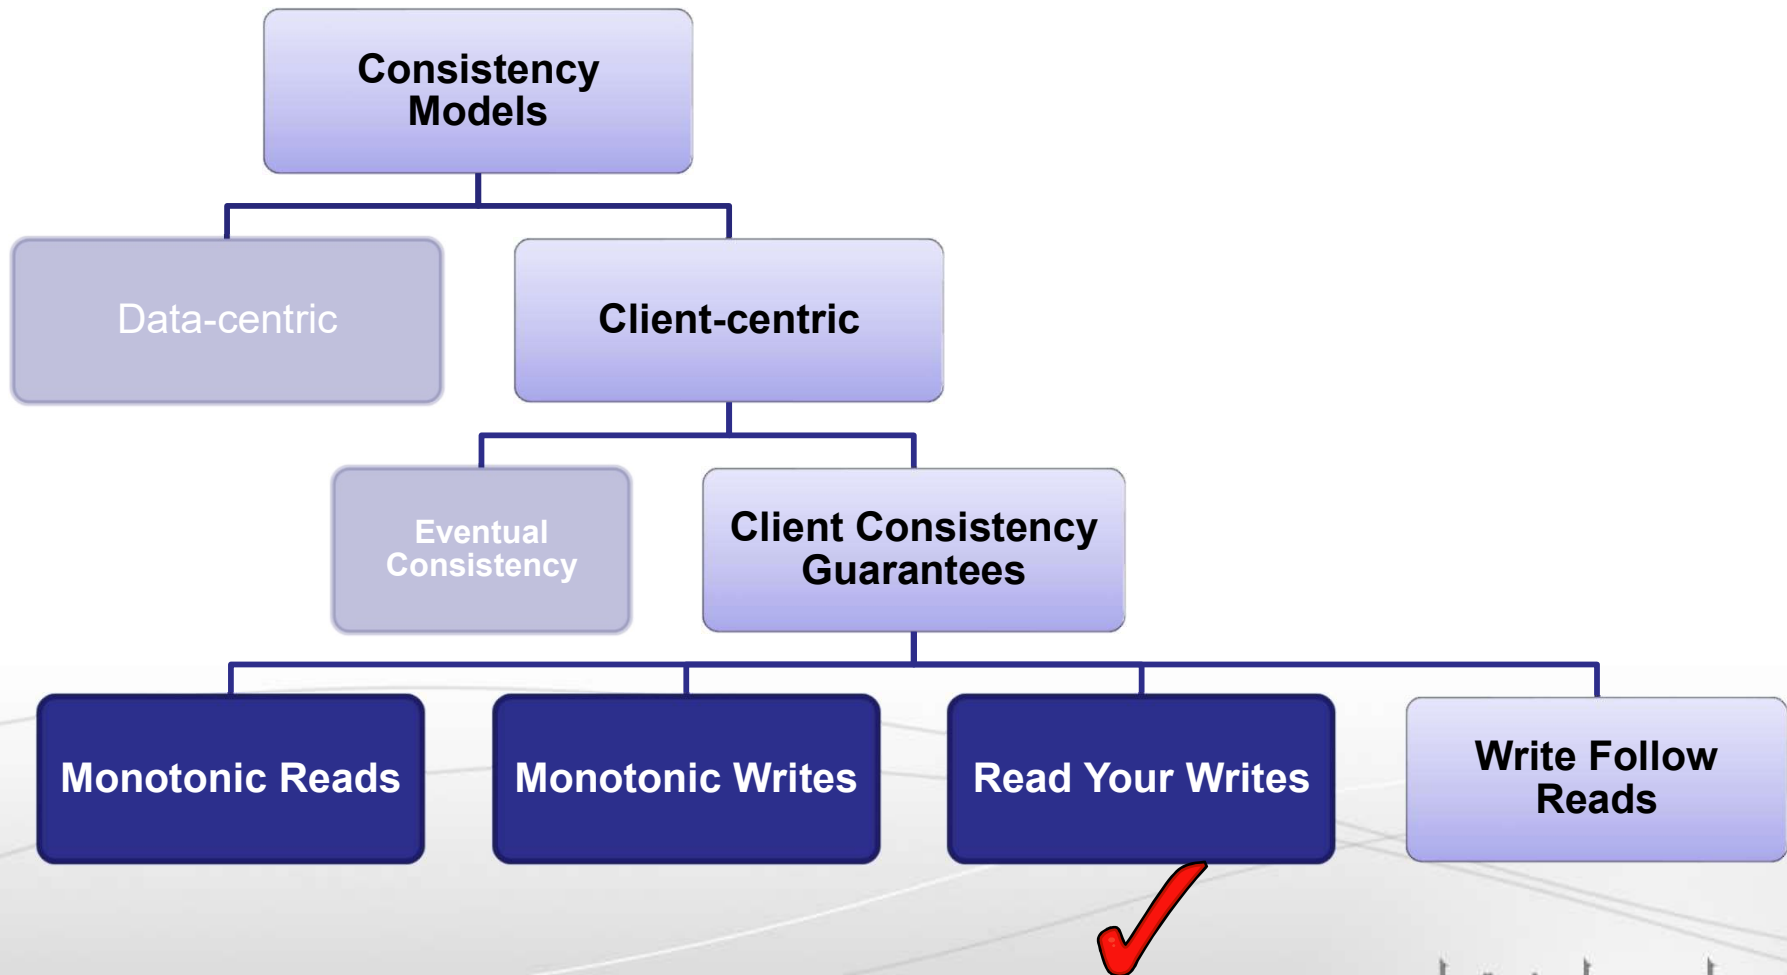
$W(x_2)$  operation should be performed only after the result of  $W(x_1)$  has been updated at  $L_2$

The data-store does not provide monotonic write consistency

# Monotonic Writes – An Example


- + Example: Updating individual libraries in a large software source code which is replicated
  - + Updates can be propagated in a lazy fashion
  - + Updates are performed on a part of the data item
    - + Some functions in an individual library is often modified and updated
  - + Monotonic writes: If an update is performed on a library, then all preceding updates on the same library are first updated
- + Question: If the update overwrites the complete software source code, is it necessary to update all the previous updates?

# Overview




# Read Your Writes

- + The effect of a write operation on a data item  $x$  by a process will always be seen by a successive read operation on  $x$  by the same process
- + Example scenario:
  - + In systems where password is stored in a replicated data-base, the password change should be seen immediately

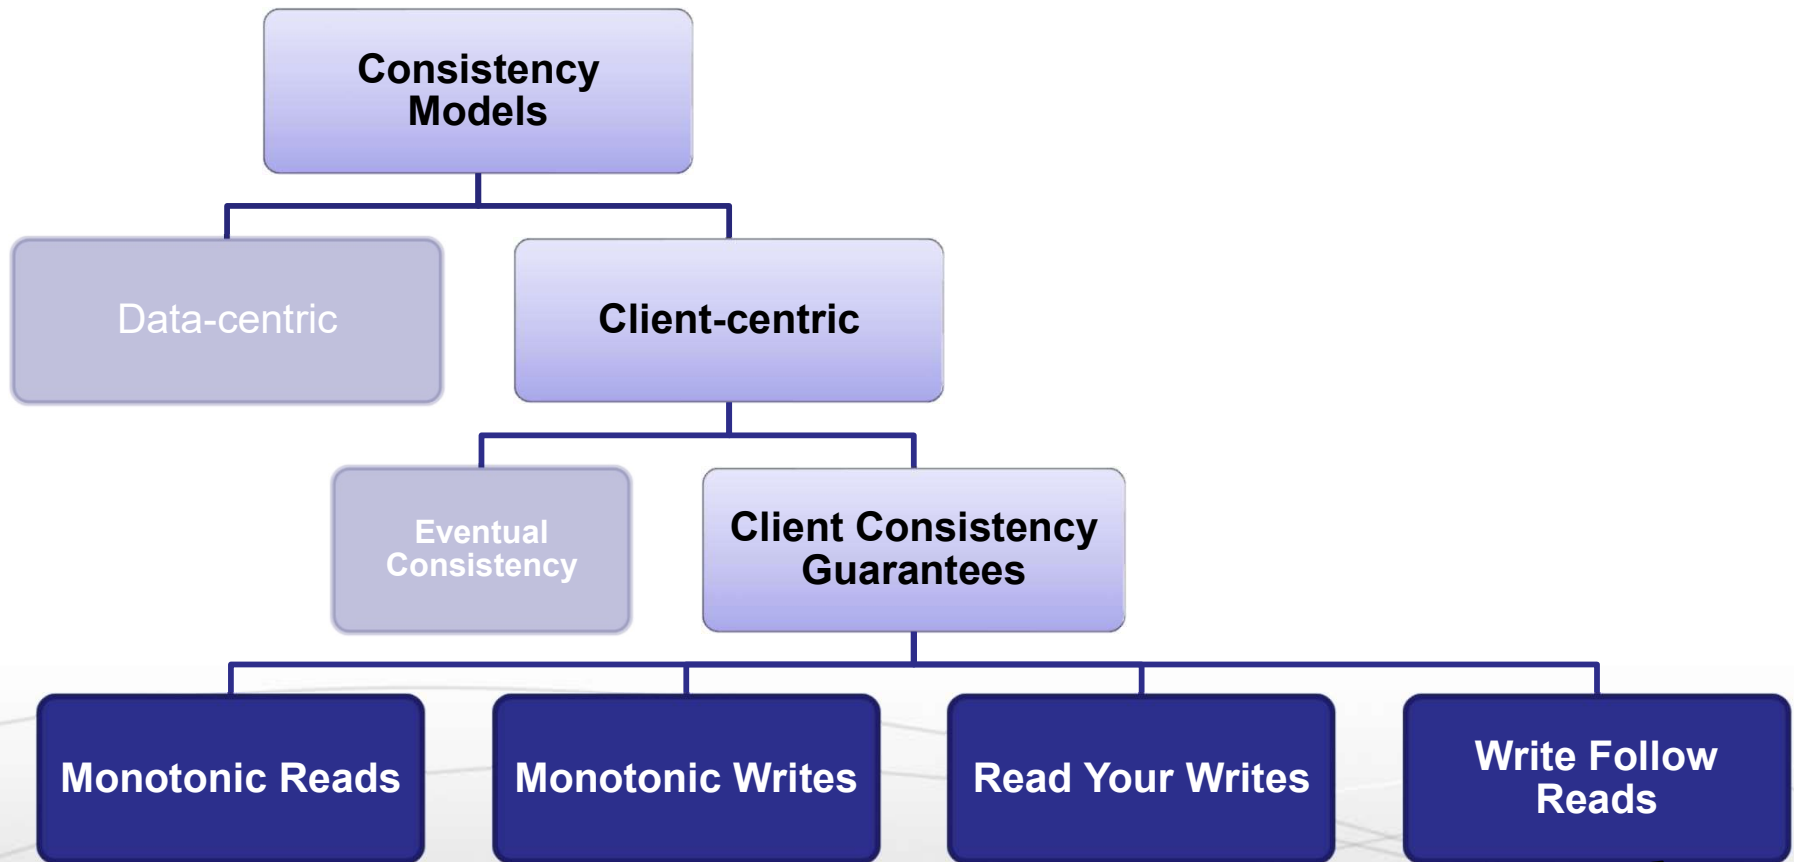


$R(x_2)$  operation should be performed only after the updating the Write Set  $WS(x_1)$  at  $L_2$



A data-store that does not provide *Read Your Write* consistency

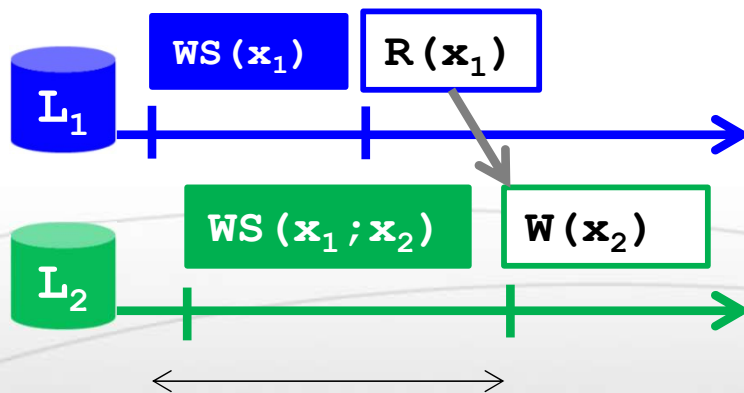
# Overview



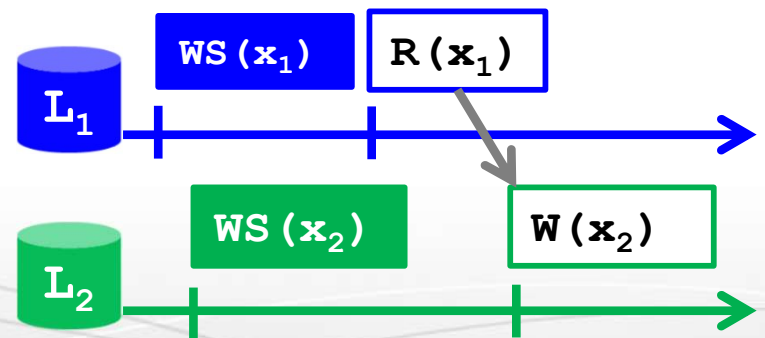


# Write Follow Reads

- + A write operation by a process on a data item  $x$  following a previous read operation on  $x$  by the same process is guaranteed to take place on the same or a more recent value of  $x$  that was read
- + Example scenario:
  - + Users of a newsgroup should post their comments only after they have read all previous comments



$W(x_2)$  operation should be performed only after the all previous writes have been seen



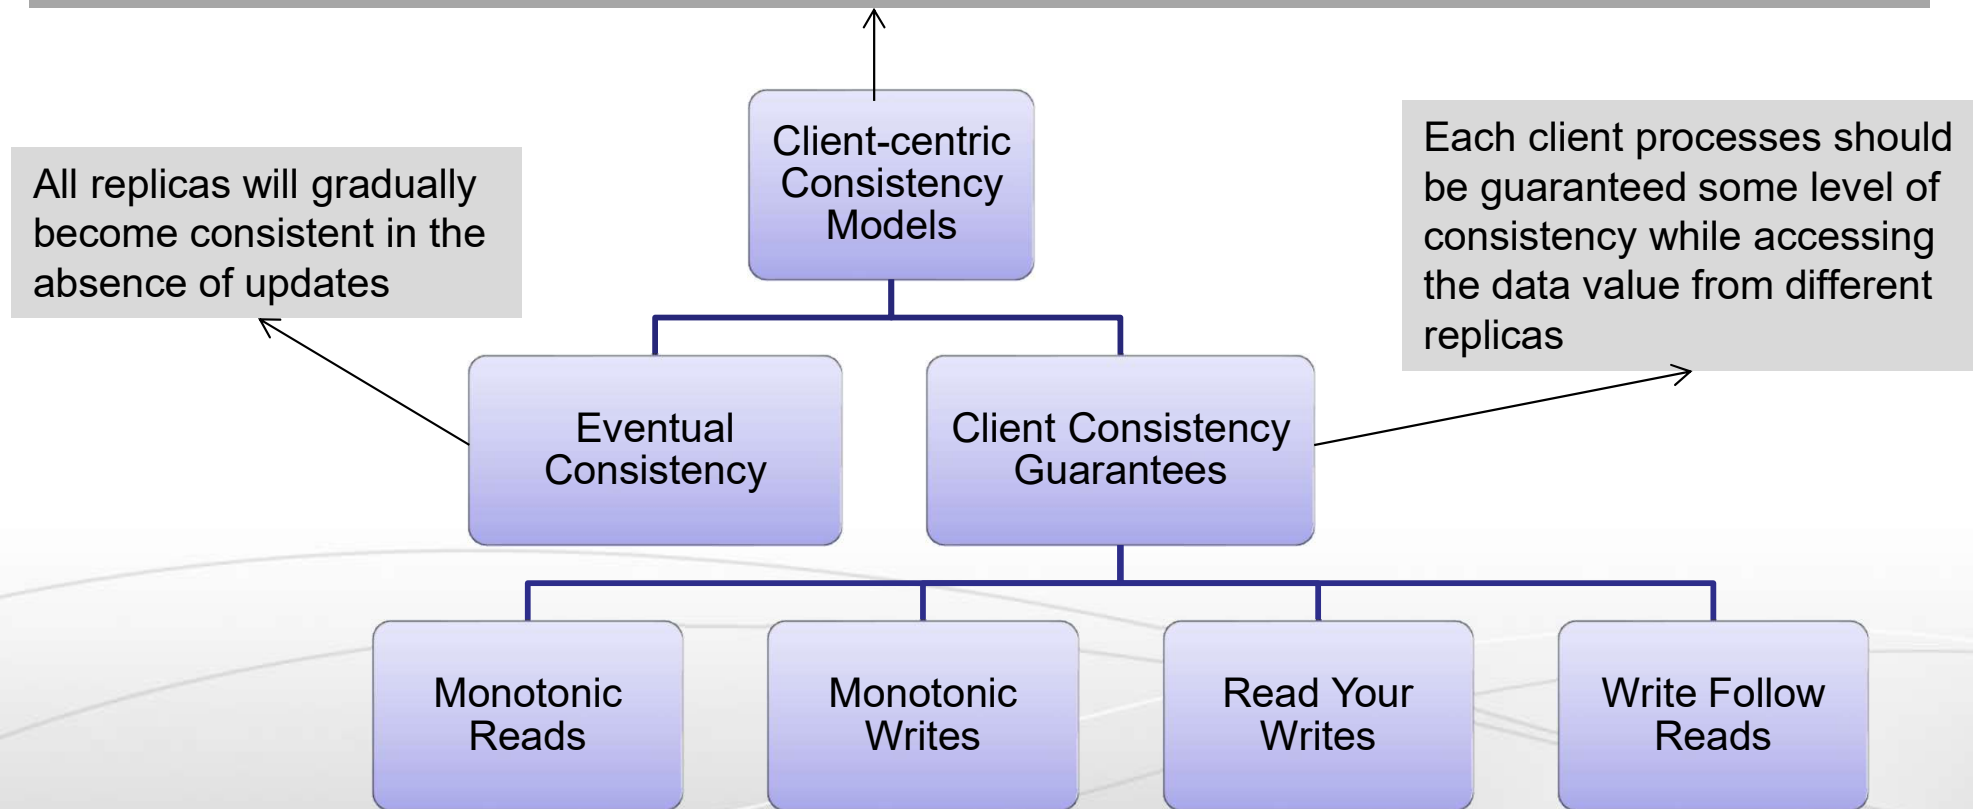
A data-store that does not guarantee Write Follow Read Consistency Model

# Summary of Client-centric Consistency Models

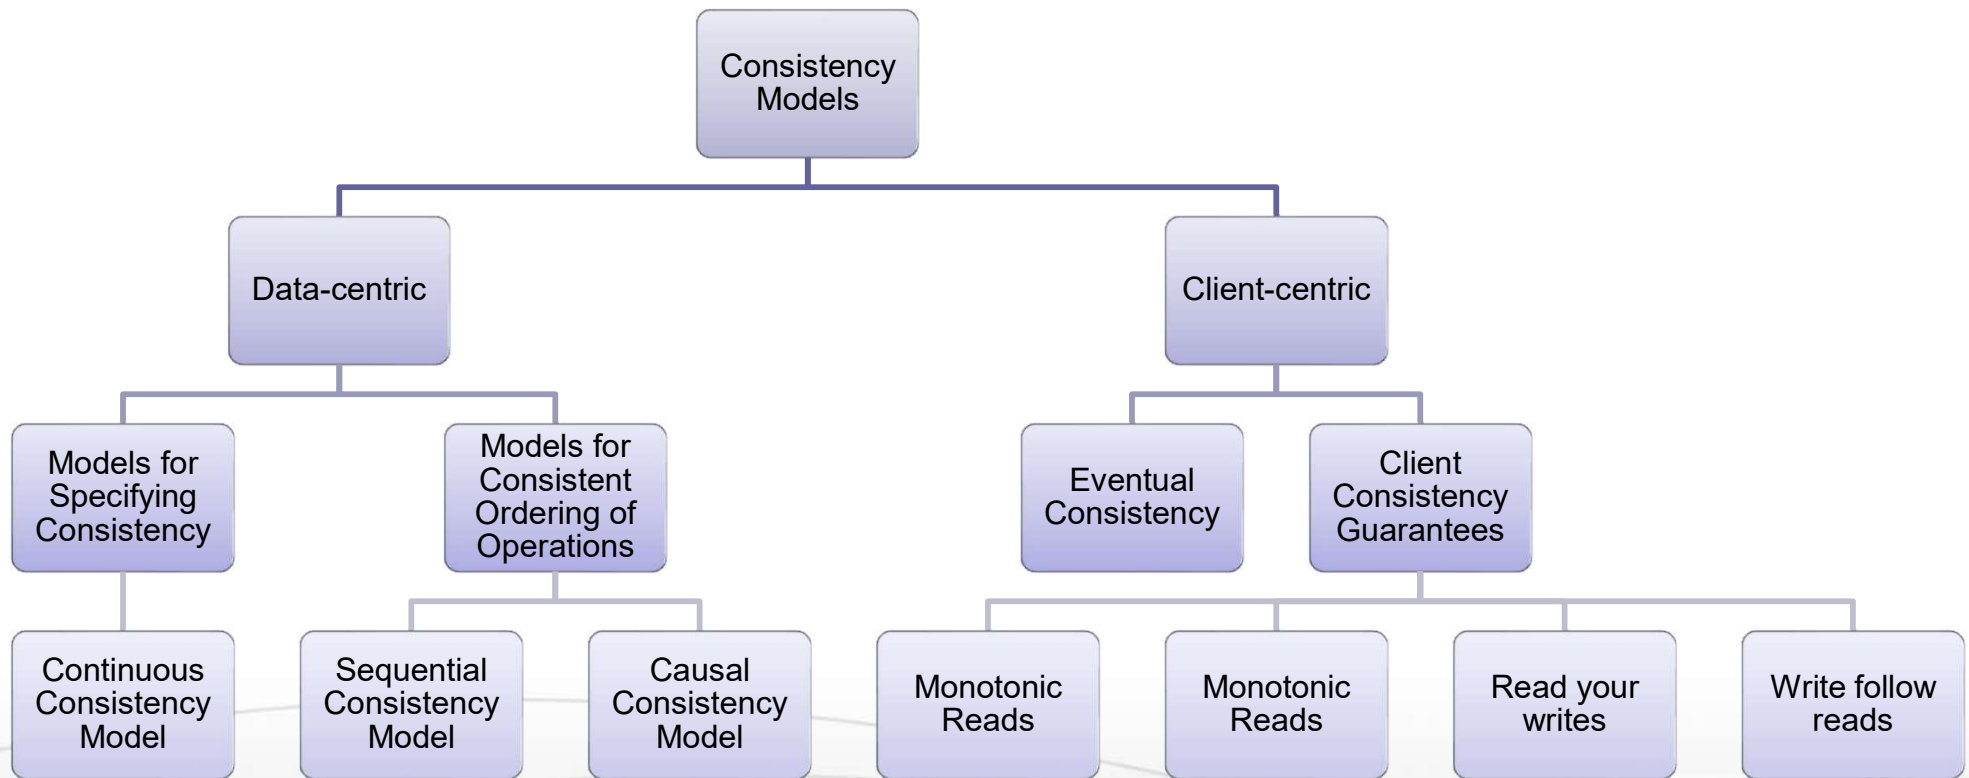
Client-centric Consistency Model defines how a data-store presents the data value to an individual client when the client process accesses the data value across different replicas.

It is generally useful in applications where:

- one client always updates the data-store.
- read-to-write ratio is high



# Topics covered in Consistency Models



# Summary of Consistency Models

- + Different applications require different levels of consistency
  - + Data-centric consistency models
    - + Define how replicas in a data-store maintain consistency
  - + Client-centric consistency models
    - + Provide an efficient, but weaker form of consistency when
    - + Here, one client process updates the data item, and many processes read the replica

# Next Class

- + Replica Management
  - + Describes where, when and by whom replicas should be placed
- + Consistency Protocols
  - + We study “how” consistency is ensured in distributed systems

# References

- + [1] Terry, D.B., Demers, A.J., Petersen, K., Spreitzer, M.J., Theimer, M.M., Welch, B.B., "Session guarantees for weakly consistent replicated data", Proceedings of the Third International Conference on Parallel and Distributed Information Systems, 1994
- + [2] Lili Qiu, Padmanabhan, V.N., Voelker, G.M., "On the placement of Web server replicas", Proceedings of IEEE INFOCOM 2001.
- + [3] Rabinovich, M., Rabinovich, I., Rajaraman, R., Aggarwal, A., "A dynamic object replication and migration protocol for an Internet hosting service", Proceedings of IEEE International Conference on Distributed Computing Systems (ICDCS), 1999
- + [4] <http://www.cdk5.net>