



Vanishing and Exploding Gradients Problems in Deep Learning

In the realm of deep learning, the optimization process plays a crucial role in training neural networks. Gradient descent, a fundamental optimization algorithm, can sometimes encounter two common issues: vanishing gradients and exploding gradients. In this article, we will delve into these challenges, providing insights into what they are, why they occur, and how to mitigate them. We will build and train a model, and learn how to face vanishing and exploding problems.

What is Vanishing Gradient?

The vanishing [gradient](#) problem is a challenge that emerges during backpropagation when the derivatives or slopes of the activation functions become progressively smaller as we move backward through the layers of a neural network. This phenomenon is particularly prominent in deep networks with many layers, hindering the effective training of the model. The weight updates becomes extremely tiny, or even exponentially small, it can significantly prolong the training time, and in the worst-case scenario, it can halt the training process altogether.

Why the Problem Occurs?

During backpropagation, the gradients propagate back through the layers of the network, they decrease significantly. This means that as they leave the output layer and return to the input layer, the gradients become progressively smaller. As a result, the weights associated with the initial levels, which accommodate these small gradients, are updated little or not at each iteration of the optimization process.

The vanishing gradient problem is particularly associated with the sigmoid and hyperbolic tangent (tanh) [activation functions](#) because their derivative



closely resemble the original ones. This persistence of small updates contributes to the vanishing gradient issue.



The sigmoid and tanh functions limit the input values to the ranges $[0,1]$ and $[-1,1]$, so that they saturate at 0 or 1 for sigmoid and -1 or 1 for Tanh. The derivatives at points becomes zero as they are moving. In these regions, especially when inputs are very small or large, the gradients are very close to zero. While this may not be a major concern in shallow networks with a few layers, it is a more pronounced issue in deep networks. When the inputs fall in saturated regions, the gradients approach zero, resulting in little update to the weights of the previous layer. In simple networks this does not pose much of a problem, but as more layers are added, these small gradients, which multiply between layers, decay significantly and consequently the first layer tears very slowly, and hinders overall model performance and can lead to convergence failure.

How can we identify?

Identifying the vanishing gradient problem typically involves monitoring the training dynamics of a deep neural network.

- One key indicator is observing model weights **converging to 0** or stagnation in the improvement of the model's performance metrics over training epochs.
- During training, if the **loss function fails to decrease** significantly, or if

- Additionally, examining the gradients themselves during backpropagation can provide insights. **Visualization techniques**, such as gradient histograms or norms, can aid in assessing the distribution of gradients throughout the network.

How can we solve the issue?

- **Batch Normalization** : Batch normalization normalizes the inputs of each layer, reducing internal covariate shift. This can help stabilize and accelerate the training process, allowing for more consistent gradient flow.
- **Activation function**: Activation function like **Rectified Linear Unit (ReLU)** can be used. With **ReLU**, the gradient is 0 for negative and zero input, and it is 1 for positive input, which helps alleviate the vanishing gradient issue. Therefore, ReLU operates by replacing poor enter values with 0, and 1 for fine enter values, it preserves the input unchanged.
- **Skip Connections and Residual Networks (ResNets)**: Skip connections, as seen in ResNets, allow the gradient to bypass certain layers during backpropagation. This facilitates the flow of information through the network, preventing gradients from vanishing.
- **Long Short-Term Memory Networks (LSTMs) and Gated Recurrent Units (GRUs)**: In the context of recurrent neural networks (RNNs), architectures like LSTMs and GRUs are designed to address the vanishing gradient problem in sequences by incorporating gating mechanisms .
- **Gradient Clipping**: Gradient clipping involves imposing a threshold on the gradients during backpropagation. Limit the magnitude of gradients during backpropagation, this can prevent them from becoming too small or exploding, which can also hinder learning.

Build and train a model for Vanishing Gradient Problem

let's see how the problems occur , and way to handle them.

Step 1: Import Libraries

First, import the necessary libraries for model

```
import numpy as np
import pandas as pd
import tensorflow as tf
import keras
from sklearn.datasets import make_moons
from sklearn.model_selection import train_test_split
from keras.layers import Dense
from keras.models import Sequential
import seaborn as sns
from imblearn.over_sampling import RandomOverSampler
from keras.layers import Dense, Dropout
from keras.optimizers import Adam
from keras.callbacks import EarlyStopping
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report
```

Step 2: Loading dataset

The code loads two CSV files (Credit_card.csv and Credit_card_label.csv) into Pandas DataFrames, df and labels.

Link to dataset: [Credit card details Binary classification.](#)

Python3

```
df = pd.read_csv('/content/Credit_card.csv')
labels = pd.read_csv('/content/Credit_card_label.csv')
```

Step 3: Data Preprocessing

We create a new column 'Approved' in the DataFrame by converting the 'label' column from the 'labels' DataFrame to integers.

Python3

```
dep = 'Approved'

df[dep] = labels.label.astype(int)

df.loc[df[dep] == 1, 'Status'] = 'Approved'
df.loc[df[dep] == 0, 'Status'] = 'Declined'
```

We perform some feature engineering on the data, creating new columns 'Age', 'EmployedDaysOnly', and 'UnemployedDaysOnly' based on existing columns.



It converts categorical variables in the 'cats' list to numerical codes using pd.Categorical and fills missing values with the mode of each column.

Python3

```
cats = [  
    'GENDER', 'Car_Owner', 'Propert_Owner', 'Type_Income',  
    'EDUCATION', 'Marital_status', 'Housing_type', 'Mobile_phone',  
    'Work_Phone', 'Phone', 'Type_Occupation', 'EMAIL_ID'  
]  
  
conts = [  
    'CHILDREN', 'Family_Members', 'Annual_income',  
    'Age', 'EmployedDaysOnly', 'UnemployedDaysOnly'  
]  
  
def proc_data():  
    df['Age'] = -df.Birthday_count // 365  
    df['EmployedDaysOnly'] = df.Employed_days.apply(lambda x: x if x > 0 else  
    df['UnemployedDaysOnly'] = df.Employed_days.apply(lambda x: abs(x) if x < 0  
  
    for cat in cats:  
        df[cat] = pd.Categorical(df[cat])  
  
    modes = df.mode().iloc[0]  
    df.fillna(modes, inplace=True)  
  
proc_data()
```

Step 5: Oversampling due to heavily skewed data and Data Splitting

Python3

```
X_over, y_over = RandomOverSampler().fit_resample(X, y)
X_train, X_val, y_train, y_val = train_test_split(X_over, y_over, test_size=0.2
```

Step 6: Encoding

The code applies the `cat.codes` method to each column specified in the `cats` list. The method is applicable to Pandas categorical data types and assigns a unique numerical code to each unique category in the categorical variable. The result is that the categorical variables are replaced with their corresponding numerical codes.

Python3

```
X_train[cats] = X_train[cats].apply(lambda x: x.cat.codes)
X_val[cats] = X_val[cats].apply(lambda x: x.cat.codes)
```

Step 7: Model Creation

Create a Sequential model using Keras. A Sequential model allows you to build a neural network by stacking layers one after another.

Python3

```
model = Sequential()
```

Step 8: Adding layers

Adding 10 dense layers to the model. Each dense layer has 10 units (neurons) and uses the sigmoid activation function. The first layer specifies `input_dim=18`, indicating that the input data has 18 features. This is the

Python3

```
model.add(Dense(10, activation='sigmoid', input_dim=18))
model.add(Dense(10, activation='sigmoid'))
model.add(Dense(10, activation='sigmoid'))
model.add(Dense(10, activation='sigmoid'))
model.add(Dense(10, activation='sigmoid'))
model.add(Dense(10, activation='sigmoid'))
model.add(Dense(10, activation='sigmoid'))
model.add(Dense(10, activation='sigmoid'))
model.add(Dense(10, activation='sigmoid'))
model.add(Dense(1, activation='sigmoid'))
```

Step 9: Model Compilation

This step specifies the loss function, optimizer, and evaluation metrics.

Python

```
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```



Step 10: Model training

Train the model using the training data (X_train and y_train) for 100 epochs. The training history is stored in the history object, which contains information about the training process, including loss and accuracy at each epoch.

Python3

```
history = model.fit(X_train, y_train, epochs=100)
```

Output:

Epoch 1/100

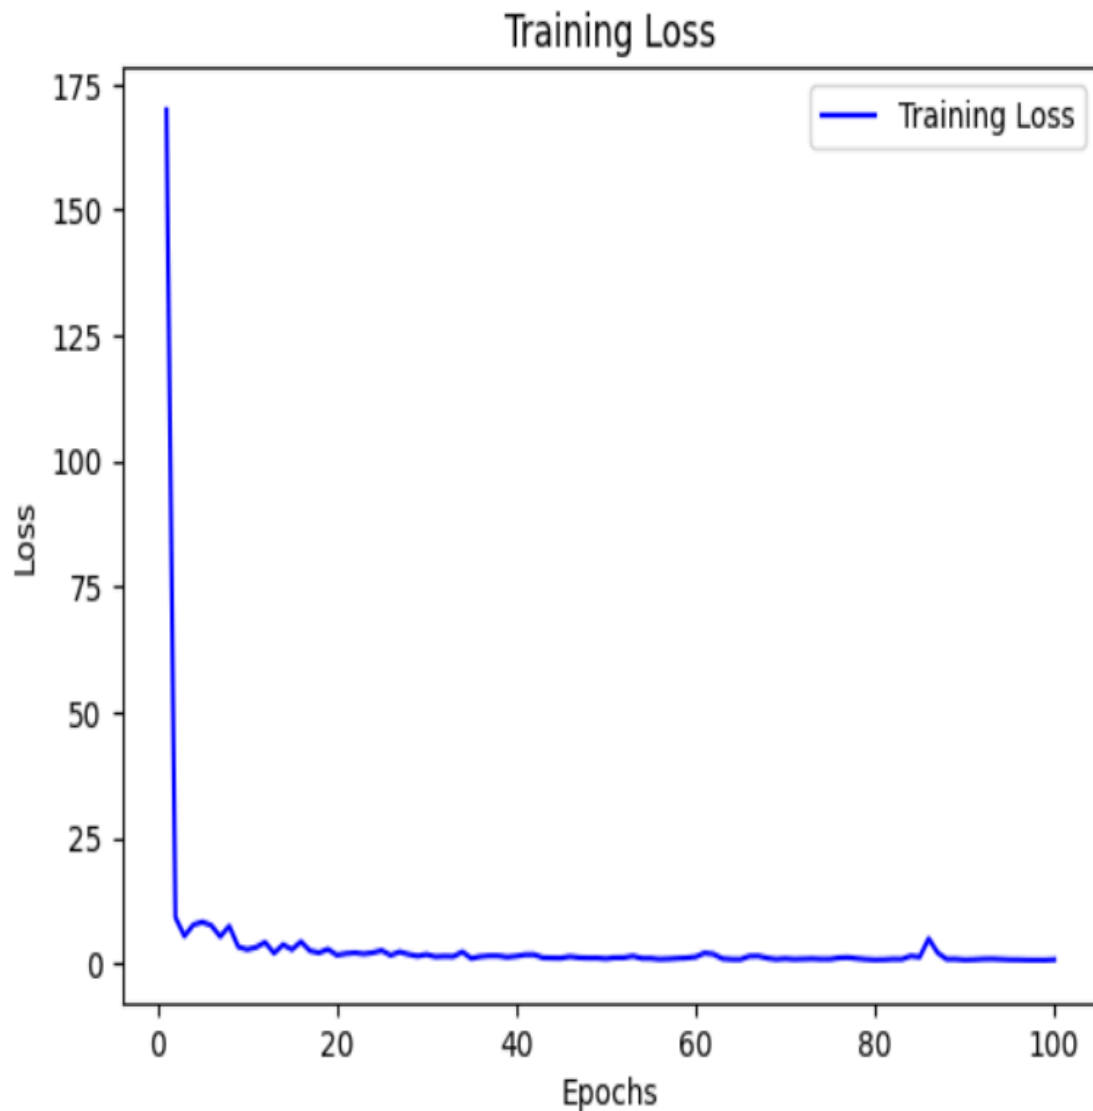
```
65/65 [=====] - 0s 3ms/step - loss: 0.6936 -  
accuracy: 0.5119  
Epoch 3/100  
65/65 [=====] - 0s 3ms/step - loss: 0.6933 -  
accuracy: 0.5119  
.  
.  
Epoch 97/100  
65/65 [=====] - 0s 3ms/step - loss: 0.6930 -  
accuracy: 0.5119  
Epoch 98/100  
65/65 [=====] - 0s 3ms/step - loss: 0.6930 -  
accuracy: 0.5119  
Epoch 99/100  
65/65 [=====] - 0s 3ms/step - loss: 0.6932 -  
accuracy: 0.5119  
Epoch 100/100  
65/65 [=====] - 0s 3ms/step - loss: 0.6929 -  
accuracy: 0.5119
```

Step 11: Plotting the training loss

Python3

```
loss = history.history['loss']  
epochs = range(1, len(loss) + 1)  
plt.plot(epochs, loss, 'b', label='Training Loss')  
plt.title('Training Loss')  
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.legend()  
plt.show()
```

Output:



Solution for Vanishing Gradient Problem

Step 1: Scaling

Python3

```
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_val_scaled = scaler.transform(X_val)
```

Step 2: Modify the Model

1. Deeper Architecture: Augment model with more layers with increased

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

2. Early Stopping: Early stopping is implemented to monitor the validation loss. Training will stop if the validation loss does not improve for a certain number of epochs (defined by patience).
3. Increased Dropout: Dropout layers are added after each dense layer to help prevent overfitting.
4. Adjusting Learning Rate: The learning rate is set to 0.001. You can experiment with different learning rates.

Python3

```
model2 = Sequential()

model2.add(Dense(128, activation='relu', input_dim=18))
model2.add(Dropout(0.5))
model2.add(Dense(256, activation='relu'))
model2.add(Dropout(0.5))
model2.add(Dense(128, activation='relu'))
model2.add(Dropout(0.5))
model2.add(Dense(64, activation='relu'))
model2.add(Dropout(0.5))
model2.add(Dense(1, activation='sigmoid'))

early_stopping = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)

model2.compile(loss='binary_crossentropy', optimizer=Adam(lr=0.001), metrics=['accuracy'])

history2 = model2.fit(X_train_scaled, y_train, epochs=100, validation_data=(X_val_scaled, y_val),
```

Output:

```
Epoch 1/100
65/65 [=====] - 3s 8ms/step - loss: 0.7167 - accuracy: 0.5308 - val_loss: 0.6851 - val_accuracy: 0.5590
Epoch 2/100
65/65 [=====] - 0s 5ms/step - loss: 0.6967 - accuracy: 0.5367 - val_loss: 0.6771 - val_accuracy: 0.6259
Epoch 3/100
65/65 [=====] - 0s 5ms/step - loss: 0.6879 -
```

```

accuracy: 0.5673 - val_loss: 0.6628 - val_accuracy: 0.6114
.
.
Epoch 96/100
65/65 [=====] - 0s 7ms/step - loss: 0.1763 -
accuracy: 0.9349 - val_loss: 0.1909 - val_accuracy: 0.9301
Epoch 97/100
65/65 [=====] - 0s 7ms/step - loss: 0.1653 -
accuracy: 0.9325 - val_loss: 0.1909 - val_accuracy: 0.9345
Epoch 98/100
65/65 [=====] - 1s 8ms/step - loss: 0.1929 -
accuracy: 0.9237 - val_loss: 0.1975 - val_accuracy: 0.9229
Epoch 99/100
65/65 [=====] - 1s 9ms/step - loss: 0.1846 -
accuracy: 0.9281 - val_loss: 0.1904 - val_accuracy: 0.9330
Epoch 100/100
65/65 [=====] - 0s 7ms/step - loss: 0.1885 -
accuracy: 0.9228 - val_loss: 0.1981 - val_accuracy: 0.9330

```

Evaluation metrics

Python3

```

predictions = model2.predict(X_val_scaled)
rounded_predictions = np.round(predictions)
report = classification_report(y_val, rounded_predictions)
print(f'Classification Report:\n{report}')

```

Output:

```
22/22 [=====] - 0s 2ms/step
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.87	0.93	352
1	0.88	1.00	0.94	335
accuracy			0.93	687
macro avg	0.94	0.93	0.93	687

What is Exploding Gradient?

The exploding gradient problem is a challenge encountered during training deep neural networks. It occurs when the gradients of the network's loss function with respect to the weights (parameters) become excessively large.

Why Exploding Gradient Occurs?

The issue of exploding gradients arises when, during backpropagation, the derivatives or slopes of the neural network's layers grow progressively larger as we move backward. This is essentially the opposite of the vanishing gradient problem.

The root cause of this problem lies in the weights of the network, rather than the choice of activation function. High weight values lead to correspondingly high derivatives, causing significant deviations in new weight values from the previous ones. As a result, the gradient fails to converge and can lead to the network oscillating around local minima, making it challenging to reach the global minimum point.

In summary, exploding gradients occur when weight values lead to excessively large derivatives, making convergence difficult and potentially preventing the neural network from effectively learning and optimizing its parameters.

As we discussed earlier, the update for the weights during backpropagation in a neural network is given by:

$$W_i = LW_i$$

Where, L is the learning rate.

The exploding gradient problem occurs when the gradients become very large during backpropagation. This is often the result of gradients greater than 1, leading to a rapid increase in values as you propagate them backward through the layers.

Mathematically, the update rule becomes problematic when $W_i > 1$, causing the weights to increase exponentially during training.

Identifying the presence of exploding gradients in deep neural network requires careful observation and analysis during training. Here are some key indicators:

- The loss function exhibits erratic behavior, oscillating wildly instead of steadily decreasing suggesting that the network weights are being updated excessively by large gradients, preventing smooth convergence.
- The training process encounters “NaN” (Not a Number) values in the loss function or other intermediate calculations..
- If network weights, during training exhibit significant and rapid increases in their values, it suggests the presence of exploding gradients.
- Tools like TensorBoard can be used to visualize the gradients flowing through the network.

How can we solve the issue?

- **Gradient Clipping:** It sets a maximum threshold for the magnitude of gradients during backpropagation. Any gradient exceeding the threshold is clipped to the threshold value, preventing it from growing unbounded.
- **Batch Normalization:** This technique normalizes the activations within each mini-batch, effectively scaling the gradients and reducing their variance. This helps prevent both vanishing and exploding gradients, improving stability and efficiency.

Build and train a model for Exploding Gradient Problem

We work on the same preprocessed data from the Vanishing gradient example but define a different neural network.

Step 1: Model creation and adding layers

Python3

```
model = Sequential()  
  
model.add(Dense(10, activation='tanh', kernel_initializer=random_normal(mean=0.  
model.add(Dense(10, activation='tanh', kernel_initializer=random_normal(mean=0.  
model.add(Dense(10, activation='tanh', kernel_initializer=random_normal(mean=0.
```

```
model.add(Dense(10, activation='tanh', kernel_initializer=random_normal(mean=0.
model.add(Dense(1, activation='sigmoid'))
# Using a poor weight initialization (random_normal with a large std deviation)
```

Step 2: Model compiling

Python3

```
optimizeroptimizer = SGD(learning_rate=1.0)
model.compile(loss='binary_crossentropy', optimizer=optimizer, metrics=['accura
```

Step 3: Model training

Python3

```
history = model.fit(X_train, y_train, epochs=100)
```


Output:

```
Epoch 1/100
65/65 [=====] - 2s 5ms/step - loss: 0.7919 -
accuracy: 0.5032
Epoch 2/100
65/65 [=====] - 0s 4ms/step - loss: 0.7440 -
accuracy: 0.5017
.
.
Epoch 99/100
65/65 [=====] - 0s 4ms/step - loss: 0.7022 -
accuracy: 0.5085
Epoch 100/100
65/65 [=====] - 0s 5ms/step - loss: 0.7037 -
accuracy: 0.5061
```

Python3

```
loss = history.history['loss'] epochs = range(1, len(loss) + 1) # Accessing the
plt.plot(epochs, loss, 'b', label='Training Loss')
plt.title('Training Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

Output:



It is observed that the loss does not converge and keeps fluctuating which shows we have encountered an exploding gradient problem.

Solution for Exploding Gradient Problem

Below methods can be used to modify the model:

1. Weight Initialization: The weight initialization is changed to 'glorot_uniform,' which is a commonly used initialization for neural networks.
2. Gradient Clipping: The clipnorm parameter in the Adam optimizer is set to 1.0, which performs gradient clipping. This helps prevent exploding gradients.

preventing exploding gradients.

Python3

```
model = Sequential()

model.add(Dense(10, activation='tanh', kernel_initializer='glorot_uniform', kernel_regularizer=keras.regularizers.l2(0.01)))
model.add(Dense(10, activation='tanh', kernel_initializer='glorot_uniform', kernel_regularizer=keras.regularizers.l2(0.01)))
model.add(Dense(10, activation='tanh', kernel_initializer='glorot_uniform', kernel_regularizer=keras.regularizers.l2(0.01)))
model.add(Dense(10, activation='tanh', kernel_initializer='glorot_uniform', kernel_regularizer=keras.regularizers.l2(0.01)))
model.add(Dense(10, activation='tanh', kernel_initializer='glorot_uniform', kernel_regularizer=keras.regularizers.l2(0.01)))
model.add(Dense(10, activation='tanh', kernel_initializer='glorot_uniform', kernel_regularizer=keras.regularizers.l2(0.01)))
model.add(Dense(10, activation='tanh', kernel_initializer='glorot_uniform', kernel_regularizer=keras.regularizers.l2(0.01)))
model.add(Dense(1, activation='sigmoid'))

early_stopping = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)

model.compile(loss='binary_crossentropy', optimizer=Adam(lr=0.001, clipnorm=1.0), metrics=['accuracy'])

history = model.fit(X_train_scaled, y_train, epochs=100, validation_data=(X_val_scaled, y_val), callbacks=[early_stopping])
```

Output:

```
Epoch 1/100
65/65 [=====] - 6s 11ms/step - loss: 0.6865 - accuracy: 0.5537 - val_loss: 0.6818 - val_accuracy: 0.5764
Epoch 2/100
65/65 [=====] - 1s 8ms/step - loss: 0.6608 - accuracy: 0.6202 - val_loss: 0.6746 - val_accuracy: 0.6070
Epoch 3/100
65/65 [=====] - 1s 8ms/step - loss: 0.6440 - accuracy: 0.6357 - val_loss: 0.6624 - val_accuracy: 0.6099
.
.
Epoch 68/100
65/65 [=====] - 1s 11ms/step - loss: 0.1909 - accuracy: 0.9257 - val_loss: 0.3819 - val_accuracy: 0.8486
Epoch 69/100
```


65/65 [=====] - 1s 10ms/step - loss: 0.1836 -
accuracy: 0.9276 - val_loss: 0.3641 - val_accuracy: 0.8515

Evaluation metrics

Python3

```
predictions = model.predict(X_val)
rounded_predictions = np.round(predictions)
report = classification_report(y_val, rounded_predictions)
print(f'Classification Report:\n{report}')
```

Output:

```
22/22 [=====] - 0s 2ms/step
Classification Report:
              precision    recall  f1-score   support

     0           0.98        0.74        0.85         352
     1           0.78        0.99        0.87         335
 accuracy                   0.86         687
 macro avg           0.88        0.86        0.86         687
 weighted avg          0.89        0.86        0.86         687
```

Conclusion

These techniques and architectural choices aim to ensure that gradients during backpropagation are within a reasonable range, enabling deep neural networks to train more effectively and converge to better solutions.

Here's a complete roadmap for you to become a developer:

Learn DSA -> Master Frontend/Backend/Full Stack -> Build Projects ->

Keep Applying to Jobs

And why go anywhere else when our **DSA to Development: Coding Guide**

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

[Development Program](#) and our counsellors will connect with you for further guidance & support.

Get paid for your published articles and stand a chance to win tablet, smartwatch and exclusive GfG goodies! Submit your entries in Dev Scriptor 2024 today.

Last Updated : 13 Dec, 2023

8

Previous

AWD-LSTM : Unraveling the Secrets of DropConnect in LSTM

Next

Python Database Optimization with Psycopg2 and Multiprocessing

Share your thoughts in the comments

Add Your Comment

Similar Reads

Deep Belief Network (DBN) in Deep Learning

Deep Boltzmann Machines (DBMs) in Deep Learning

Unveiling the Power of Fastai: A Deep Dive into the Versatile Deep Learning Library

Higher-Order gradients in TensorFlow

Longformer in Deep Learning

Kaiming Initialization in Deep Learning

Recursive Neural Network in Deep Learning

Top 10 best frameworks for deep learning in 2024

Top 10 Deep Learning Tools You Must Know [2024]

List of Deep Learning Layers

Complete Tutorials

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

Brain Teasers

SDLC Models | Software Development Models

UPSC Notification 2024: Exam Date, Eligibility, Syllabus

P paulsubhro

Article Tags : [Geeks Premier League 2023](#) , [Picked](#) , [Deep Learning](#) , [Geeks Premier League](#)

Additional Information



“My educational aspiration has finally met its passion in our country's best university.”

Now I can proudly say **#IHaveMadeIt**

With SBI Scholar Education Loan

- ✓ Collateral free loan up to Rs. 50 lakhs
- ✓ Nil Processing Fee
- ✓ Up to 100% Finance

Dreaming of pursuing studies India's leading universities? Let's make it happen.

 State Bank of India

 **GeeksforGeeks**
Sanchhaya Education Private Limited
A-143, 9th Floor, Sovereign Corporate Tower, Sector-136, Noida, Uttar Pradesh - 201305



[Company](#)

[Explore](#)

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

Careers
In Media
Contact Us
Advertise with us
GFG Corporate Solution
Placement Training Program

GfG Weekly Contest
Offline Classes (Delhi/NCR)
DSA in JAVA/C++
Master System Design
Master CP
GeeksforGeeks Videos
Geeks Community

Languages

Python
Java
C++
PHP
GoLang
SQL
R Language
Android Tutorial
Tutorials Archive

Data Science & ML

Data Science With Python
Data Science For Beginner
Machine Learning Tutorial
ML Maths
Data Visualisation Tutorial
Pandas Tutorial
NumPy Tutorial
NLP Tutorial
Deep Learning Tutorial

Python

Python Programming Examples
Django Tutorial
Python Projects
Python Tkinter

DSA

Data Structures
Algorithms
DSA for Beginners
Basic DSA Problems
DSA Roadmap
Top 100 DSA Interview Problems
DSA Roadmap by Sandeep Jain
All Cheat Sheets

HTML & CSS

HTML
CSS
Web Templates
CSS Frameworks
Bootstrap
Tailwind CSS
SASS
LESS
Web Design

Computer Science

GATE CS Notes
Operating Systems
Computer Network

Database Management System

DevOps

Git
AWS
Docker
Kubernetes
Azure
GCP
DevOps Roadmap

System Design

High Level Design
Low Level Design
UML Diagrams
Interview Guide
Design Patterns
OOAD
System Design Bootcamp
Interview Questions

NCERT Solutions

Class 12
Class 11
Class 10
Class 9
Class 8
Complete Study Material

Commerce

Accountancy
Business Studies
Economics
Management
HR Management

Competitive Programming

Top DS or Algo for CP
Top 50 Tree
Top 50 Graph
Top 50 Array
Top 50 String
Top 50 DP
Top 15 Websites for CP

JavaScript

JavaScript Examples
TypeScript
ReactJS
NextJS
AngularJS
NodeJS
Lodash
Web Browser

School Subjects

Mathematics
Physics
Chemistry
Biology
Social Science
English Grammar

UPSC Study Material

Polity Notes
Geography Notes
History Notes
Science and Technology Notes
Economy Notes

SSC/ BANKING

SSC CGL Syllabus
SBI PO Syllabus
SBI Clerk Syllabus
IBPS PO Syllabus
IBPS Clerk Syllabus
SSC CGL Practice Papers

Companies

META Owned Companies
Alphabet Owned Companies
TATA Group Owned Companies
Reliance Owned Companies
Fintech Companies
EdTech Companies

Exams

JEE Mains
JEE Advanced
GATE CS
NEET
UGC NET

Free Online Tools

Typing Test
Image Editor
Code Formatters
Code Converters
Currency Converter
Random Number Generator
Random Password Generator

Colleges

Indian Colleges Admission & Campus Experiences
List of Central Universities - In India
Colleges in Delhi University
IIT Colleges
NIT Colleges
IIIT Colleges

Preparation Corner

Company-Wise Recruitment Process
Resume Templates
Aptitude Preparation
Puzzles
Company-Wise Preparation

More Tutorials

Software Development
Software Testing
Product Management
SAP
SEO - Search Engine Optimization
Linux
Excel

Write & Earn

Write an Article
Improve an Article
Pick Topics to Write
Share your Experiences
Internships